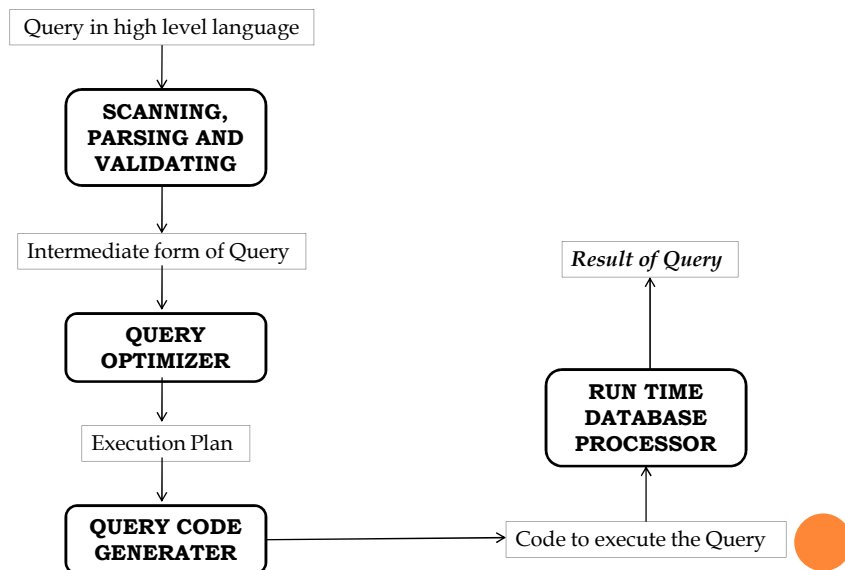


# Query Optimization

Swarup Kr Ghosh

## Introduction



## Query Optimization

- ✓ **What is Query Optimization ?**
- ✓ **What is the utility of Query Optimization?**
- ✓ **Query Optimization is based on what ?**

❖ *Relational Algebra*

(Select ( $\sigma$ ), Project ( $\pi$ ), Cartesian Product ( $\times$ ), Join ( $\bowtie$ ) etc)

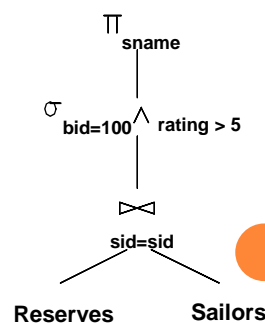


## QUERY OPTIMIZATION OVERVIEW

- Query can be converted to relational algebra
- Rel. Algebra converted to tree, joins as branches
- Each operator has implementation choices
- Operators can also be applied in different order!

```
SELECT S.sname
FROM Reserves R, Sailors S
WHERE R.sid=S.sid AND
      R.bid=100 AND S.rating>5
```

$\pi_{(sname)} \sigma_{(bid=100 \wedge rating > 5)} (Reserves \bowtie Sailors)$



# Query Optimization Technique

- ✓ **Heuristics Technique in Query Optimization.**
- ✓ **Cost Estimates Technique in Query Optimization.**

## Heuristics Technique of Query Optimization

- One of the main heuristic rules is to apply **SELECT** and **PROJECT** operations before applying the **JOIN** or other binary operations.
- Generating equivalent Query Tree and Query Graph.

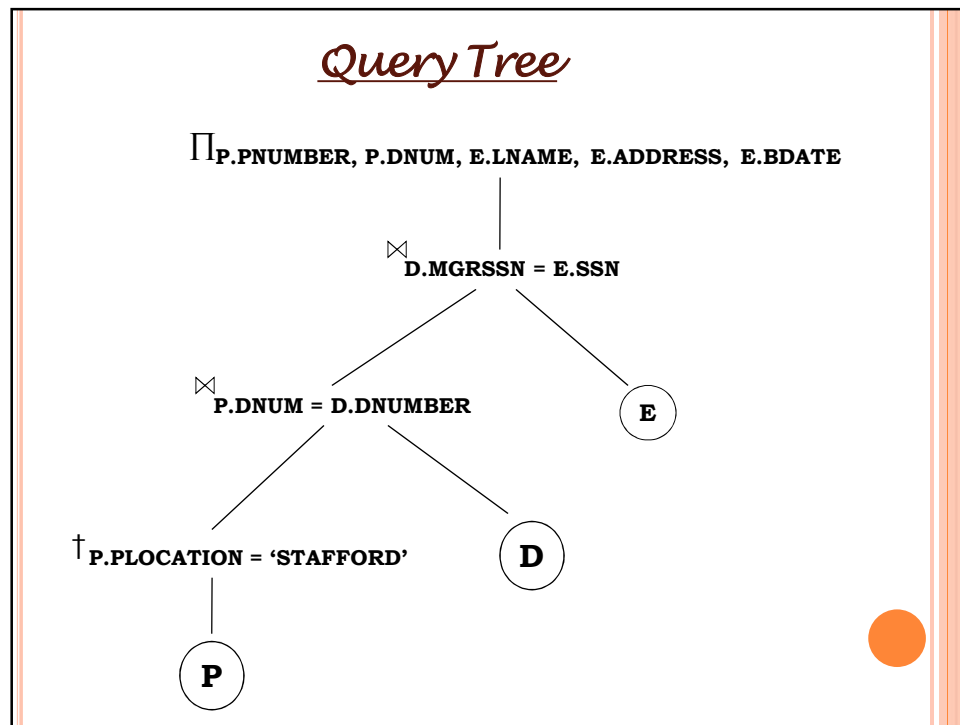
Example:

SQL Query: **SELECT P.PNUMBER, D.DNUM, E.LNAME, E.ADDRESS, E.BDATE  
FROM PROJECT AS P, DEPARTMENT AS D, EMPLOYEE AS E  
WHERE D.MGRSSN = E.SSN AND P.DNUM = D.DNUMBER  
AND P.PLOCATION = 'STAFFORD'**

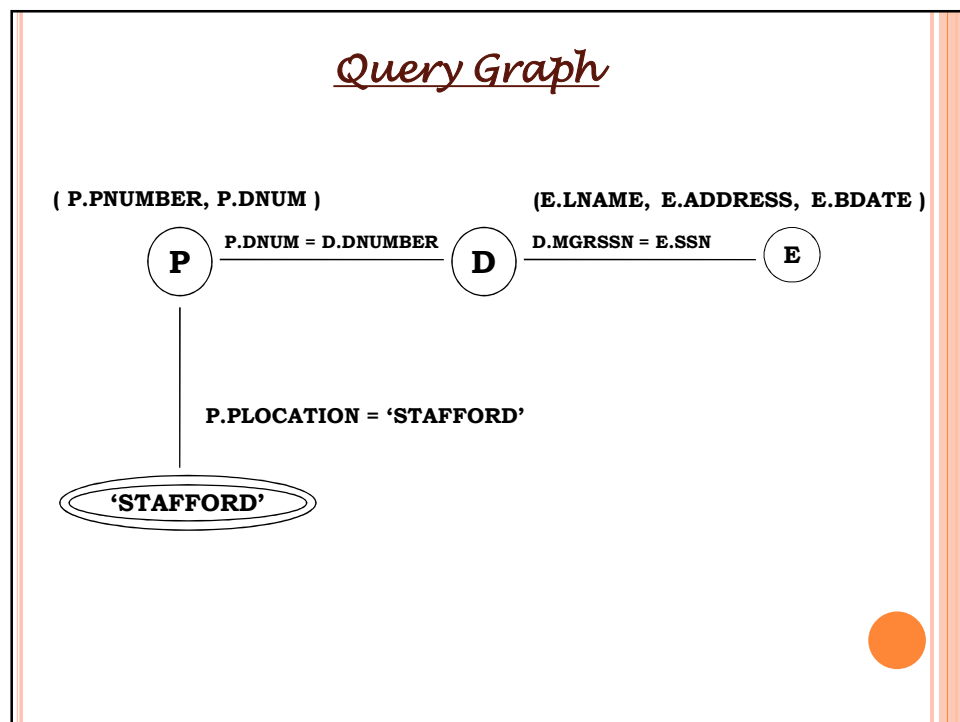
Relational Algebra:

$\pi_{PNUMBER, DNUM, LNAME, ADDRESS, BDATE} (\sigma_{PLOCATION = 'STAFFORD'} (PROJECT) \bowtie \sigma_{DNUM = DNUMBER} (DEPARTMENT) \bowtie \sigma_{MGRSSN = SSN} (EMPLOYEE))$

## Query Tree



## Query Graph



## General Transformation Rules for Relational Algebra Operations

1. **Cascade of  $\sigma$**  : A conjunctive selection condition can be broken up into a cascade (that is a sequence) of individual  $\sigma$  operations:

$$\sigma_{c1 \text{ and } c2 \text{ and } \dots \text{ and } cn} (R) \equiv \sigma_{c1}(\sigma_{c2}(\dots(\sigma_{cn}(R))\dots))$$

2. **Commutative of  $\sigma$** : The  $\sigma$  operation is commutative.

$$\sigma_{c1}(\sigma_{c2}(R)) \equiv \sigma_{c2}(\sigma_{c1}(R))$$

3. **Cascade of  $\Pi$**  : In a cascade (sequence) of  $\Pi$  operations, all but the last one can be ignored:

$$\Pi_{list1}(\Pi_{list2}(\dots(\Pi_{listn}(R))\dots)) \equiv \Pi_{list1}(R)$$

4. **Commuting  $\sigma$  with  $\Pi$** : If the selection condition  $c$  involves only those attributes  $A1, \dots, An$  in the projection list, the two operations can be commuted.

$$\Pi_{A1, A2, \dots, An}(\sigma_c(R)) \equiv \sigma_c(\Pi_{A1, A2, \dots, An}(R))$$

## General Transformation Rules for Relational Algebra Operations (Cont.)

5. **Commutative of  $\bowtie$  (and  $X$ )** : The  $\bowtie$  operation is commutative, as is the  $X$  operations:

$$R \bowtie_c S \equiv S \bowtie_c R$$

$$R X S \equiv S X R$$

6. **Commutative  $\sigma$  with  $\bowtie$  (or  $X$ )**: If all the attributes in the selection condition  $c$  involve only the attributes of one of the relations being joined – say,  $R$  – the two operations can be commuted as follows:

$$\sigma_c(R \bowtie S) \equiv (\sigma_c(R)) \bowtie S$$

Alternatively, if the selection condition  $c$  can be written as  $(c1 \text{ AND } c2)$ , where condition  $c1$  involves only the attributes of  $R$  and condition  $c2$  involves only the attributes of  $S$ , the operations commute as follows:

$$\sigma_c(R \bowtie S) \equiv (\sigma_{c1}(R)) \bowtie (\sigma_{c2}(S))$$

The same rules apply if the  $\bowtie$  is replaced by a  $X$  operation

### General Transformation Rules for Relational Algebra Operations (Cont.)

**7. Commuting  $\Pi$  with  $\bowtie$  (or  $X$ ):** Suppose that the projection list is  $L = \{ A_1, \dots, A_n, B_1, \dots, B_m \}$ , where  $A_1, \dots, A_n$  are attributes of  $R$  and  $B_1, \dots, B_m$  are attributes of  $S$ . If the join condition  $c$  involves only attributes in  $L$ , the two operations can be commuted as follows:

$$\Pi_L(R \bowtie_c S) \equiv (\Pi_{A_1, \dots, A_n}(R)) \bowtie_c (\Pi_{B_1, \dots, B_m}(S))$$

If the join condition  $c$  contains additional attributes not in  $L$ , these must be added to the projection list, and a final  $\Pi$  operation is needed. For example, if attributes  $A_{n+1}, \dots, A_{n+k}$  of  $R$  and  $B_{m+1}, \dots, B_{m+p}$  of  $S$  are involved in the join condition  $c$  but are not in the projection list  $L$ , the operations commute as follows:

$$\Pi_L(R \bowtie_c S) \equiv \Pi_L(\Pi_{A_1, \dots, A_n, A_{n+1}, \dots, A_{n+k}}(R)) \bowtie_c (\Pi_{B_1, \dots, B_m, B_{m+1}, \dots, B_{m+p}}(S)).$$

For  $X$ , there is no condition  $c$ , so the first transformation rule always applies by replacing  $c$  with  $X$ .

**8. Commutativity of set operations:** The set operations  $\cap$  and  $\cup$  are commutative but  $-$  is not.

### General Transformation Rules for Relational Algebra Operations (Cont.)

**9. Associativity of  $\bowtie$ ,  $X$ ,  $\cup$ , and  $\cap$ :** These four operations are individually associative; that is, if  $\theta$  stands for any one of these four operations (throughout the expression), we have:

$$(R \theta S) \theta T \equiv R \theta (S \theta T)$$

**10. Commuting  $\sigma$  with set operations:** The  $\sigma$  operation commutes with  $\cup$ ,  $\cap$ , and  $-$ . If  $\theta$  stands for any one of these three operations (throughout the expression), we have:

$$\sigma_c(R \theta S) \equiv (\sigma_c(R)) \theta (\sigma_c(S))$$

**11. The  $\Pi$  operation commutes with  $\cup$ :**

$$\Pi_L(R \cup S) \equiv (\Pi_L(R)) \cup (\Pi_L(S))$$

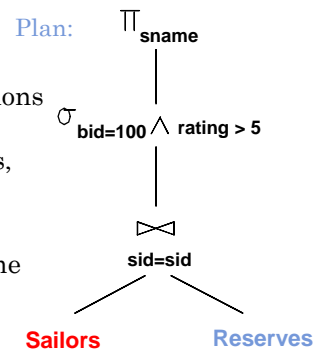
**12. Converting a  $(\sigma, X)$  sequence into  $\bowtie$ :** If the condition  $c$  of a  $\sigma$  that follows a  $X$  corresponds to a join condition, convert the  $(\sigma, X)$  sequence into a  $\bowtie$  as follows:

$$(\sigma_c(R X S)) \equiv (R \bowtie_c S)$$

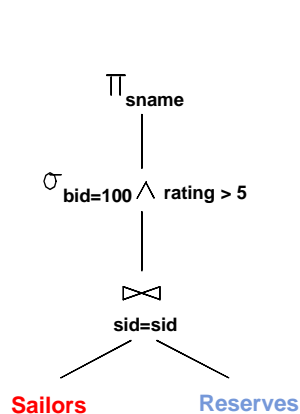
## MOTIVATING EXAMPLE

```
SELECT S.sname
FROM Reserves R, Sailors S
WHERE R.sid=S.sid AND
      R.bid=100 AND S.rating>5
```

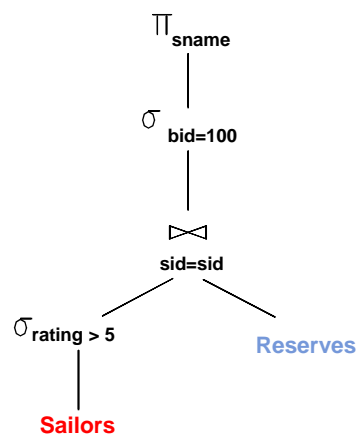
- Cost: ?
- By no means the worst plan!
- Misses several opportunities: selections could have been 'pushed' earlier, no use is made of any available indexes, etc.
- Goal of optimization:* To find more efficient plans that compute the same answer.



## ALTERNATIVE PLANS

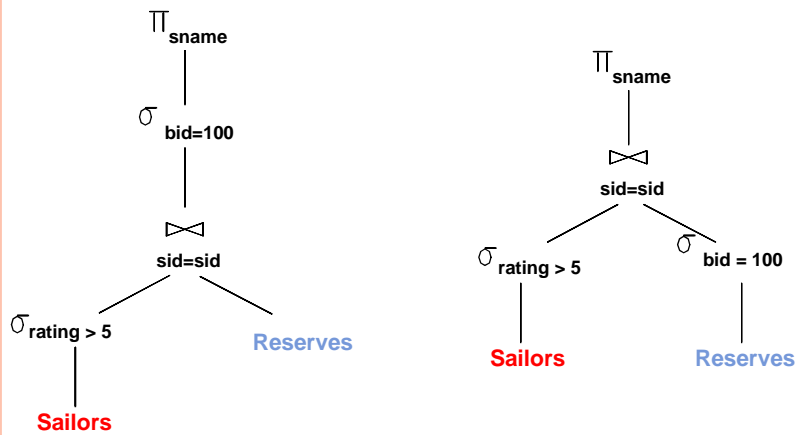


More operations



Less operations

## ALTERNATIVE PLANS



## Example :

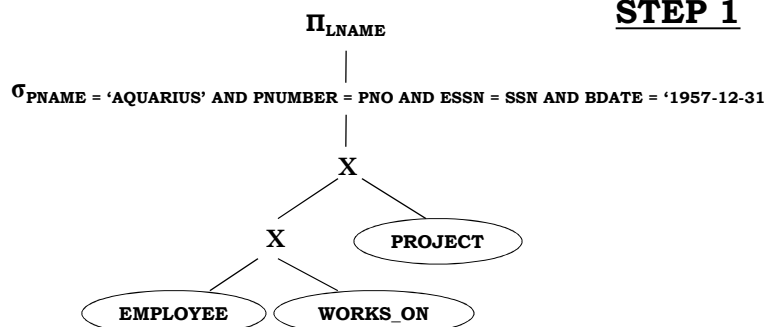
### SQL Query:

```
SELECT LNAME
FROM EMPLOYEE, WORKS_ON, PROJECT
WHERE PNAME = 'AQUARIUS' AND PNUMBER = PNO AND
      ESSN = SSN AND BDATE = '1957-12-31';
```

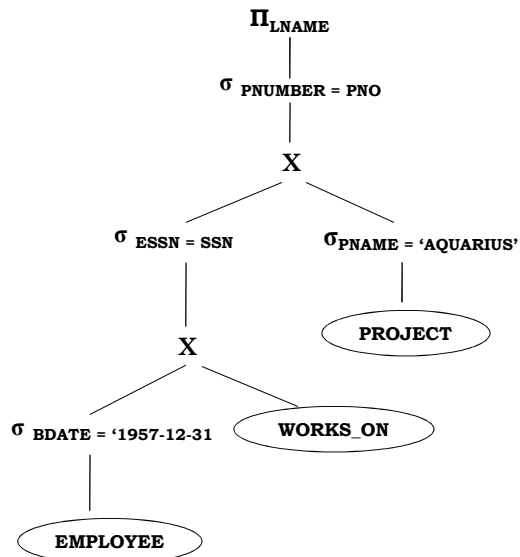
### Relational Algebra:

$\Pi_{LNAME} (\sigma_{PNAME = 'AQUARIUS' \text{ AND } PNUMBER = PNO \text{ AND } ESSN = SSN \text{ AND } BDATE = '1957-12-31'} (EMPLOYEE \times WORKS\_ON \times PROJECT))$

### STEP 1





**STEP 2****Using Rules:****Rule 1:**

$$\sigma_{c1 \text{ and } \dots \text{ and } cn}(R) \equiv \sigma_{c1}(\sigma_{c2}(\dots(\sigma_{cn}(R)) \dots))$$

**Rule 2:**

$$\sigma_{c1}(\sigma_{c2}(R)) \equiv \sigma_{c2}(\sigma_{c1}(R))$$

**Rule 4:**

$$\Pi_{A1, \dots, An}(\sigma_c(R)) \equiv \sigma_c(\Pi_{A1, \dots, An}(R))$$

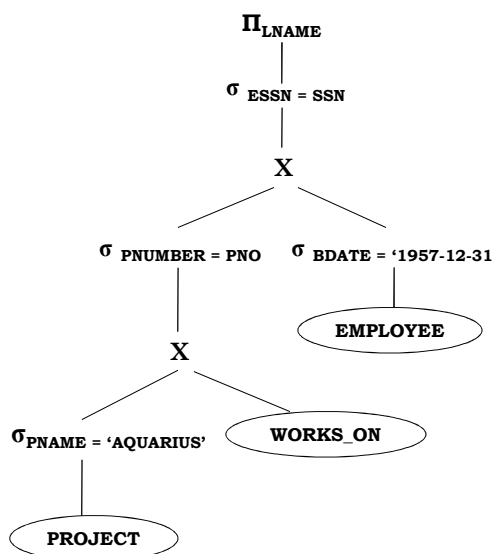
**Rule 6:**

$$\sigma_c(R \bowtie S) \equiv (\sigma_c(R)) \bowtie S$$

$$\sigma_c(R \bowtie S) \equiv (\sigma_{c1}(R)) \bowtie (\sigma_{c2}(S))$$

**Rule 10:**

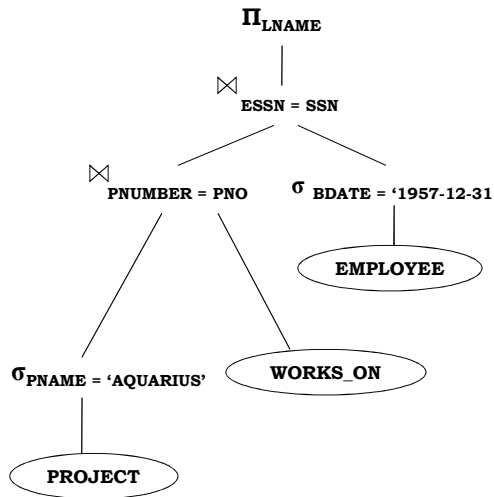
$$\sigma_c(R \theta S) \equiv (\sigma_c(R)) \theta (\sigma_c(S))$$

**STEP 3****Using Rules:****Rule 5:**

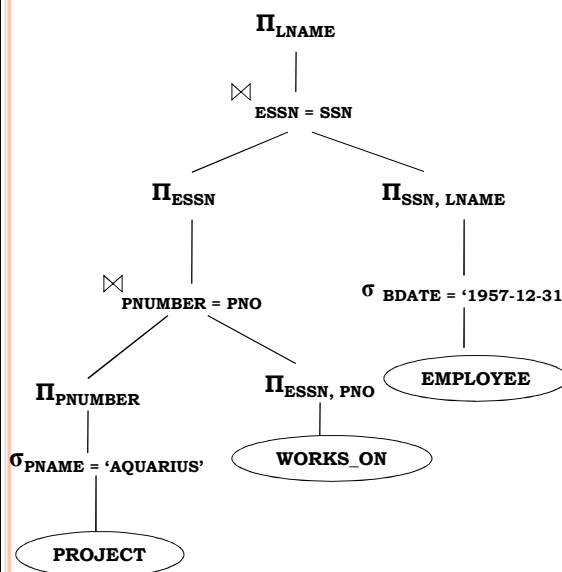
$$R \bowtie S \equiv S \bowtie R$$

**Rule 9:**

$$(R \theta S) \theta T \equiv R \theta (S \theta T)$$

**STEP 4****Using Rules:****Rule 12:**

$$(\sigma_c (R \bowtie S)) \equiv (R \bowtie_c S)$$

**STEP 5****Using Rules:****Rule 3:**

$$\Pi_{list1}(\Pi_{list2}(\dots(\Pi_{listn}(R))\dots)) \equiv \Pi_{list1}(R)$$

**Rule 4:**

$$\Pi_{A1, \dots, An}(\sigma_c(R)) \equiv \sigma_c(\Pi_{A1, \dots, An}(R))$$

**Rule 7:**

$$\Pi_L(R \bowtie_c S) \equiv (\Pi_{A1, \dots, An}(R)) \bowtie_c (\Pi_{B1, \dots, Bm}(S))$$

**Rule 11:**

$$\Pi_L(R \cup S) \equiv (\Pi_L(R)) \cup (\Pi_L(S))$$




# Selectivity and Cost Estimates in Query Optimization

## Cost Component for Query Execution

- ✓ Access cost to secondary storage
- ✓ Storage Cost
- ✓ Computation Cost
- ✓ Memory usage Cost
- ✓ Communication Cost

**NB:** Generally we use Access Cost to Secondary Storage mainly and use Communication Cost for Distributed Database only.

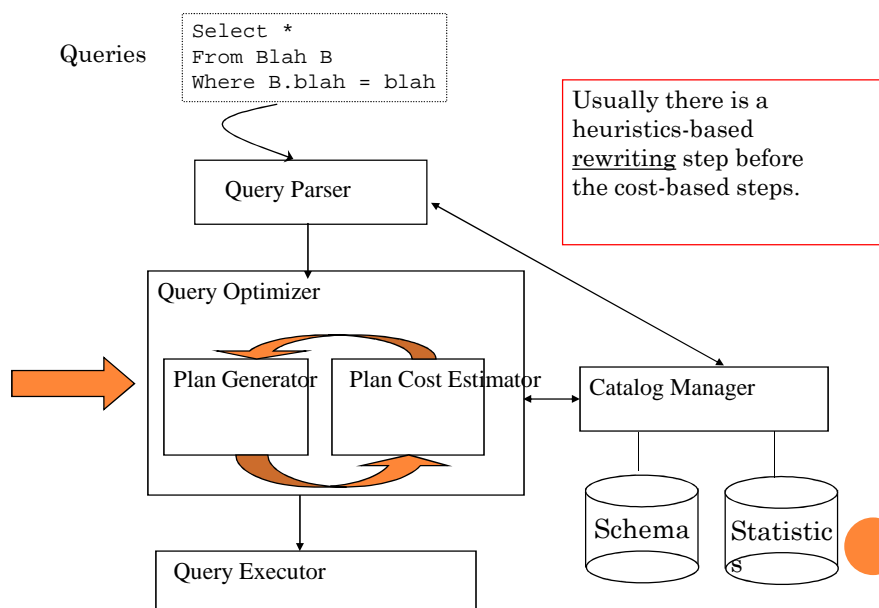


## Catalog Information Used in Cost Function

- Number of Records (tuples) ( $r$ ).
- Record Size ( $R$ )
- Number of Blocks ( $b$ )
- Blocking factor ( $bfr$ )
- Number of Levels ( $x$ )
- Number of first-level index blocks ( $b_{11}$ )
- Number of Distinct values ( $d$ )
- Selectivity ( $sl$ )
- Selection Cardinality ( $s = sl * r$ )

**NB:** For a key attribute,  $d = r$ ,  $sl = 1/r$  and  $s = 1$ . For a nonkey attribute, by making an assumption that the  $d$  distinct values are uniformly distributed among the records, we estimate  $sl = (1/d)$  and so  $s = (r/d)$

## COST-BASED QUERY SUB-SYSTEM



## SCHEMA FOR EXAMPLES

Sailors (sid: integer, sname: string, rating: integer, age: real)

Reserves (sid: integer, bid: integer, day: dates, rname: string)

- Reserves:
  - Each tuple is 40 bytes long, 100 tuples per page, 1000 pages.
  - Assume there are 100 boats
- Sailors:
  - Each tuple is 50 bytes long, 80 tuples per page, 500 pages.
  - Assume there are 10 different ratings

## Cost Functions for SELECT

- **Linear Search (Brute force) Approach.**

$$C_{s1a} = b$$

$$C_{s1b} = (b/2)$$

- **Binary Search.**

$$C_{s2} = \log_2 b + (s/bfr) - 1$$

- **Using a primary index or hash key to retrieve a single record.**

$$C_{s3a} = x + 1 \text{ (for Primary index)}$$

$$C_{s3b} = 1 \text{ (for hash key)}$$

- **Using an ordering index to retrieve multiple records.**

$$C_{s4} = x + (b/2)$$

- **Using a clustering index to retrieve multiple records.**

$$C_{s5} = x + (s/bfr)$$

### Example to Illustrate Cost Based Query Optimization

**SQL Query:** **SELECT P.PNUMBER, D.DNUM, E.LNAME, E.ADDRESS, E.BDATE**  
**FROM PROJECT AS P, DEPARTMENT AS D, EMPLOYEE AS E**  
**WHERE D.MGRSSN = E.SSN AND P.DNUM = D.DNUMBER**  
**AND P.PLOCATION = 'STAFFORD'**

TABLE_NAME	COLUMN_NAME	NUM_DISTINCT	LOW_VALUE	HIGH_VALUE
PROJECT	PLOCATION	200	1	200
PROJECT	PNUMBER	2000	1	2000
PROJECT	DNUM	50	1	50
DEPARTMENT	DNUMBER	50	1	50
DEPARTMENT	MGRSSN	50	1	50
EMPLOYEE	SSN	10000	1	10000
EMPLOYEE	DNO	50	1	50
EMPLOYEE	SALARY	500	1	500

### Example to Illustrate Cost Based Query Optimization (Cont.)

TABLE_NAME	NUM_ROWS	BLOCKS
PROJECT	2000	100
DEPARTMENT	50	5
EMPLOYEE	10000	2000

INDEX_NAME	UNIQUENES	BLEVEL	LEAF_BLOCKS	DISTINCT_KEYS
PROJ_PLOC	NONUNIQUE	1	4	200
EMP_SSN	UNIQUE	1	50	10000
EMP_SAL	NONUNIQUE	1	50	500

PROJECT ✕ DEPARTMENT ✕ EMPLOYEE  
 DEPARTMENT ✕ PROJECT ✕ EMPLOYEE  
 DEPARTMENT ✕ EMPLOYEE ✕ PROJECT  
 EMPLOYEE ✕ DEPARTMENT ✕ PROJECT

