

Introduction to SQL

@SKG

2/18/2020 1

Outline

- Overview of The SQL Query Language
- Data Definition
- Basic Query Structure
- Additional Basic Operations
- Set Operations
- Null Values
- Aggregate Functions
- Nested Subqueries
- Modification of the Database

@SKG

2/18/2020 2

History

- IBM **Sequel** language developed as part of System R project at the IBM San Jose Research Laboratory
- Renamed Structured Query Language (SQL)
- ANSI and ISO standard SQL:
 - SQL-86
 - SQL-89
 - SQL-92
 - SQL:1999 (language name became Y2K compliant!)
 - SQL:2003
- Commercial systems offer most, if not all, SQL-92 features, plus varying feature sets from later standards and special proprietary features.
 - Not all examples here may work on your particular system.

@SKG

2/18/2020 3

Data Definition Language

The SQL data-definition language (DDL) allows the specification of information about relations, including:

- The schema for each relation.
- The domain of values associated with each attribute.
- Integrity constraints
- And as we will see later, also other information such as
 - The set of indices to be maintained for each relations.
 - Security and authorization information for each relation.
 - The physical storage structure of each relation on disk.

@SKG

2/18/2020 4

Domain Types in SQL

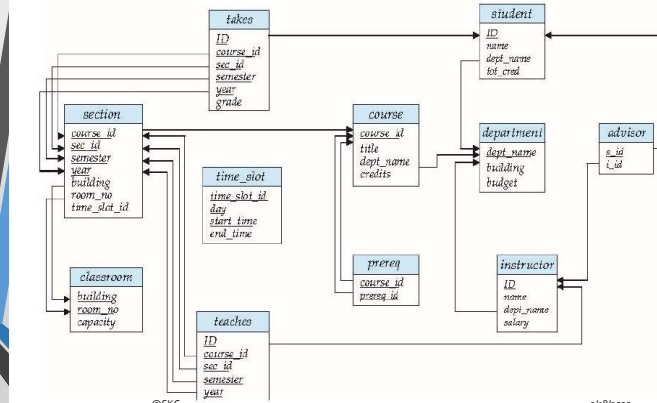
- **char(n)**. Fixed length character string, with user-specified length n .
- **varchar(n)**. Variable length character strings, with user-specified maximum length n .
- **int**. Integer (a finite subset of the integers that is machine-dependent).
- **smallint**. Small integer (a machine-dependent subset of the integer domain type).
- **numeric(p,d)**. Fixed point number, with user-specified precision of p digits, with d digits to the right of decimal point. (ex., **numeric(3,1)**, allows 44.5 to be stored exactly, but not 444.5 or 0.32)
- **real, double precision**. Floating point and double-precision floating point numbers, with machine-dependent precision.
- **float(n)**. Floating point number, with user-specified precision of at least n digits.

@SKG

2/18/2020

5

Schema Diagram for University Database



@SKG

2/18/2020

6

Create Table Construct

- An SQL relation is defined using the **create table** command:

```
create table r (A1 D1, A2 D2, ..., An Dn,
               (integrity-constraint1),
               ...,
               (integrity-constraintk))
```

- r is the name of the relation
- each A_i is an attribute name in the schema of relation r
- D_i is the data type of values in the domain of attribute A_i

- Example:

```
create table instructor (
    ID      char(5),
    name    varchar2(20),
    dept_name varchar2(20),
    salary  numeric(8,2))
```

@SKG

2/18/2020

7

Integrity Constraints in Create Table

- **not null**
- **primary key** (A_1, \dots, A_n)
- **foreign key** (A_m, \dots, A_n) **references** r

Example:

```
create table instructor (
    ID      char(5),
    name    varchar2(20) not null,
    dept_name varchar2(20),
    salary  numeric(8,2),
    primary key (ID),
    foreign key (dept_name) references department);
```

primary key declaration on an attribute automatically ensures **not null**

@SKG


2/18/2020

8

- **create table student** (
 ID varchar2(5),
 name varchar2(20) not null,
 dept_name varchar2(20),
 tot_cred numeric(3,0),
 primary key (ID),
 foreign key (dept_name) **references** department);
- **create table takes** (
 ID varchar2(5),
 course_id varchar2(8),
 sec_id varchar2(8),
 semester varchar2(6),
 year numeric(4,0),
 grade varchar2(2),
 primary key (ID, course_id, sec_id, semester, year),
 foreign key (ID) **references** student,
 foreign key (course_id, sec_id, semester, year) **references** section);
- Note: sec_id can be dropped from primary key above, to ensure a student cannot be registered for two sections of the same course in the same semester.

[illegible]

- ### Schema Diagram for University Database



Updates to tables

- **Insert**
 - insert into *instructor* values ('10211', 'Smith', 'Biology', 66000);
- **Delete**
 - Remove all tuples from the *student* relation
 - delete from *student*
- **Drop Table**
 - drop table *r*
- **Alter**
 - alter table *r* add *A D*
 - where *A* is the name of the attribute to be added to relation *r* and *D* is the domain of *A*.
 - All exiting tuples in the relation are assigned *null* as the value for the new attribute.
 - alter table *r* drop *A*
 - where *A* is the name of an attribute of relation *r*
 - Dropping of attributes not supported by many databases.

@SKG

2/18/2020 11

- @SKG

2/18/2020 11

Basic Query Structure

- A typical SQL query has the form:

```
select  $A_1, A_2, \dots, A_n$   
from  $r_1, r_2, \dots, r_m$   
where  $P$ 
```
- A_i represents an attribute
- R_i represents a relation
- P is a predicate.

The result of an SQL query is a relation.

@SKG 2/18/2020 12

- ```

select A_1, A_2, \dots, A_n
from r_1, r_2, \dots, r_m
where P

```

- @SKG

2/18/2020 12

## The select Clause

- The **select** clause lists the attributes desired in the result of a query
  - corresponds to the projection operation of the relational algebra
- Example: find the names of all instructors:  

```
select name
from instructor
```
- NOTE: SQL names are case insensitive (i.e., you may use upper- or lower-case letters.)
  - E.g.,  $Name \equiv NAME \equiv name$
  - Some people use upper case wherever we use **bold font**.

2/18/2020 13

## The select Clause (Cont.)

- SQL allows duplicates in relations as well as in query results.
- To force the elimination of duplicates, insert the keyword **distinct** after select.
- Find the department names of all instructors, and remove duplicates  

```
select distinct dept_name
from instructor
```
- The keyword **all** specifies that duplicates should not be removed.

```
select all dept_name
from instructor
```

@SKG

2/18/2020 14

## The select Clause (Cont.)

- An asterisk in the select clause denotes "all attributes"  

```
select *
from instructor
```
- An attribute can be a literal with no **from** clause  

```
select '437'
```

  - Results is a table with one column and a single row with value "437"
  - Can give the column a name using:  

```
select '437' as FOO
```
- An attribute can be a literal with **from** clause  

```
select 'A'
from instructor
```

  - Result is a table with one column and  $N$  rows (number of tuples in the *instructors* table), each row with value "A"

2/18/2020 15

## The select Clause (Cont.)

- The **select** clause can contain arithmetic expressions involving the operation,  $+$ ,  $-$ ,  $*$ , and  $/$ , and operating on constants or attributes of tuples.
  - The query:  

```
select ID, name, salary/12
from instructor
```

would return a relation that is the same as the *instructor* relation, except that the value of the attribute *salary* is divided by 12.
  - Can rename "*salary/12*" using the **as** clause:  

```
select ID, name, salary/12 as monthly_salary
```

@SKG

2/18/2020 16

## The where Clause

- The **where** clause specifies conditions that the result must satisfy
  - Corresponds to the selection predicate of the relational algebra.
- To find all instructors in Comp. Sci. dept

```
select name
from instructor
where dept_name = 'Comp. Sci.'
```

- Comparison results can be combined using the logical connectives **and**, **or**, and **not**
  - To find all instructors in Comp. Sci. dept with salary > 80000

```
select name
from instructor
where dept_name = 'Comp. Sci.' and salary > 80000
```

Comparisons can be applied to results of arithmetic expressions.

@SKG

2/18/2020 17

## The from Clause

- The **from** clause lists the relations involved in the query
  - Corresponds to the Cartesian product operation of the relational algebra.

- Find the **Cartesian product** *instructor* X *teaches*

```
select *
from instructor, teaches
```

- generates every possible instructor – teaches pair, with all attributes from both relations.
- For common attributes (e.g., *ID*), the attributes in the resulting table are renamed using the relation name (e.g., *instructor.ID*)
- Cartesian product not very useful directly, but useful combined with where-clause condition (selection operation in relational algebra).

@SKG

2/18/2020 18

## Cartesian Product

*instructor*

| ID    | name       | dept_name  | salary |
|-------|------------|------------|--------|
| 10101 | Srinivasan | Comp. Sci. | 65000  |
| 12121 | Wu         | Finance    | 90000  |
| 15151 | Mozart     | Music      | 40000  |
| 22222 | Einstein   | Physics    | 95000  |
| 32343 | El Said    | History    | 60000  |

*teaches*

| ID    | course_id | sec_id | semester | year |
|-------|-----------|--------|----------|------|
| 10101 | CS-101    | 1      | Fall     | 2009 |
| 10101 | CS-315    | 1      | Spring   | 2010 |
| 10101 | CS-347    | 1      | Fall     | 2009 |
| 12121 | FIN-201   | 1      | Spring   | 2010 |
| 15151 | MU-199    | 1      | Spring   | 2010 |
| 22222 | PHY-101   | 1      | Fall     | 2009 |

| Inst.ID | name       | dept_name  | salary | teaches.ID | course_id | sec_id | semester | year |
|---------|------------|------------|--------|------------|-----------|--------|----------|------|
| 10101   | Srinivasan | Comp. Sci. | 65000  | 10101      | CS-101    | 1      | Fall     | 2009 |
| 10101   | Srinivasan | Comp. Sci. | 65000  | 10101      | CS-315    | 1      | Spring   | 2010 |
| 10101   | Srinivasan | Comp. Sci. | 65000  | 10101      | CS-347    | 1      | Fall     | 2009 |
| 10101   | Srinivasan | Comp. Sci. | 65000  | 12121      | FIN-201   | 1      | Spring   | 2010 |
| 10101   | Srinivasan | Comp. Sci. | 65000  | 15151      | MU-199    | 1      | Spring   | 2010 |
| 10101   | Srinivasan | Comp. Sci. | 65000  | 22222      | PHY-101   | 1      | Fall     | 2009 |
| ...     | ...        | ...        | ...    | ...        | ...       | ...    | ...      | ...  |
| ...     | ...        | ...        | ...    | ...        | ...       | ...    | ...      | ...  |
| 12121   | Wu         | Finance    | 90000  | 10101      | CS-101    | 1      | Fall     | 2009 |
| 12121   | Wu         | Finance    | 90000  | 10101      | CS-315    | 1      | Spring   | 2010 |
| 12121   | Wu         | Finance    | 90000  | 10101      | CS-347    | 1      | Fall     | 2009 |
| 12121   | Wu         | Finance    | 90000  | 12121      | FIN-201   | 1      | Spring   | 2010 |
| 12121   | Wu         | Finance    | 90000  | 15151      | MU-199    | 1      | Spring   | 2010 |
| 12121   | Wu         | Finance    | 90000  | 22222      | PHY-101   | 1      | Fall     | 2009 |
| ...     | ...        | ...        | ...    | ...        | ...       | ...    | ...      | ...  |
| ...     | ...        | ...        | ...    | ...        | ...       | ...    | ...      | ...  |

19

## Examples

- Find the names of all instructors who have taught some course and the course\_id
  - select name, course\_id**  
**from instructor, teaches**  
**where instructor.ID = teaches.ID**
- Find the names of all instructors in the Art department who have taught some course and the course\_id
  - select name, course\_id**  
**from instructor, teaches**  
**where instructor.ID = teaches.ID and instructor.dept\_name = 'Art'**

@SKG

2/18/2020 20

## Cartesian Product

*instructor*

| ID    | name       | dept_name  | salary |
|-------|------------|------------|--------|
| 10101 | Srinivasan | Comp. Sci. | 65000  |
| 12121 | Wu         | Finance    | 90000  |
| 15151 | Mozart     | Music      | 40000  |
| 22222 | Einstein   | Physics    | 95000  |
| 32343 | El Said    | History    | 60000  |

*teaches*

| ID    | course_id | sec_id | semester | year |
|-------|-----------|--------|----------|------|
| 10101 | CS-101    | 1      | Fall     | 2009 |
| 10101 | CS-315    | 1      | Spring   | 2010 |
| 10101 | CS-347    | 1      | Fall     | 2009 |
| 12121 | FIN-201   | 1      | Spring   | 2010 |
| 15151 | MU-199    | 1      | Spring   | 2010 |
| 22222 | PHY-101   | 1      | Fall     | 2009 |

| Inst.ID | name       | dept_name  | salary | teaches.ID | course_id | sec_id | semester | year |
|---------|------------|------------|--------|------------|-----------|--------|----------|------|
| 10101   | Srinivasan | Comp. Sci. | 65000  | 10101      | CS-101    | 1      | Fall     | 2009 |
| 10101   | Srinivasan | Comp. Sci. | 65000  | 10101      | CS-315    | 1      | Spring   | 2010 |
| 10101   | Srinivasan | Comp. Sci. | 65000  | 10101      | CS-347    | 1      | Fall     | 2009 |
| 10101   | Srinivasan | Comp. Sci. | 65000  | 12121      | FIN-201   | 1      | Spring   | 2010 |
| 10101   | Srinivasan | Comp. Sci. | 65000  | 15151      | MU-199    | 1      | Spring   | 2010 |
| 10101   | Srinivasan | Comp. Sci. | 65000  | 22222      | PHY-101   | 1      | Fall     | 2009 |
| ...     | ...        | ...        | ...    | ...        | ...       | ...    | ...      | ...  |
| ...     | ...        | ...        | ...    | ...        | ...       | ...    | ...      | ...  |
| 12121   | Wu         | Finance    | 90000  | 10101      | CS-101    | 1      | Fall     | 2009 |
| 12121   | Wu         | Finance    | 90000  | 10101      | CS-315    | 1      | Spring   | 2010 |
| 12121   | Wu         | Finance    | 90000  | 10101      | CS-347    | 1      | Fall     | 2009 |
| 12121   | Wu         | Finance    | 90000  | 12121      | FIN-201   | 1      | Spring   | 2010 |
| 12121   | Wu         | Finance    | 90000  | 15151      | MU-199    | 1      | Spring   | 2010 |
| 12121   | Wu         | Finance    | 90000  | 22222      | PHY-101   | 1      | Fall     | 2009 |
| ...     | ...        | ...        | ...    | ...        | ...       | ...    | ...      | ...  |
| ...     | ...        | ...        | ...    | ...        | ...       | ...    | ...      | ...  |

22

## Examples

- Find the names of all instructors who have taught some course and course\_id

```
select name, course_id
from instructor, teaches
where instructor.ID = teaches.ID
```

- Equi-Join, Natural Join

| ID    | name       | dept_name  | salary |
|-------|------------|------------|--------|
| 10101 | Srinivasan | Comp. Sci. | 65000  |
| 12121 | Wu         | Finance    | 90000  |
| 15151 | Mozart     | Music      | 40000  |
| 22222 | Einstein   | Physics    | 95000  |
| 32343 | El Said    | History    | 60000  |

| Inst.ID | name       | dept_name  | salary | teaches.ID | course_id | sec_id | semester | year |
|---------|------------|------------|--------|------------|-----------|--------|----------|------|
| 10101   | Srinivasan | Comp. Sci. | 65000  | 10101      | CS-101    | 1      | Fall     | 2009 |
| 10101   | Srinivasan | Comp. Sci. | 65000  | 10101      | CS-315    | 1      | Spring   | 2010 |
| 10101   | Srinivasan | Comp. Sci. | 65000  | 10101      | CS-347    | 1      | Fall     | 2009 |
| 10101   | Srinivasan | Comp. Sci. | 65000  | 12121      | FIN-201   | 1      | Spring   | 2010 |
| 10101   | Srinivasan | Comp. Sci. | 65000  | 15151      | MU-199    | 1      | Spring   | 2010 |
| 10101   | Srinivasan | Comp. Sci. | 65000  | 22222      | PHY-101   | 1      | Fall     | 2009 |
| ...     | ...        | ...        | ...    | ...        | ...       | ...    | ...      | ...  |
| ...     | ...        | ...        | ...    | ...        | ...       | ...    | ...      | ...  |
| 12121   | Wu         | Finance    | 90000  | 10101      | CS-101    | 1      | Fall     | 2009 |
| 12121   | Wu         | Finance    | 90000  | 10101      | CS-315    | 1      | Spring   | 2010 |
| 12121   | Wu         | Finance    | 90000  | 10101      | CS-347    | 1      | Fall     | 2009 |
| 12121   | Wu         | Finance    | 90000  | 12121      | FIN-201   | 1      | Spring   | 2010 |
| 12121   | Wu         | Finance    | 90000  | 15151      | MU-199    | 1      | Spring   | 2010 |
| 12121   | Wu         | Finance    | 90000  | 22222      | PHY-101   | 1      | Fall     | 2009 |
| ...     | ...        | ...        | ...    | ...        | ...       | ...    | ...      | ...  |
| ...     | ...        | ...        | ...    | ...        | ...       | ...    | ...      | ...  |

@SKG

22

## The Rename Operation

- The SQL allows renaming relations and attributes using the **as** clause:

*old-name as new-name*

- Find the names of all instructors who have a higher salary than some instructor in 'Comp. Sci'.

```
select distinct T.name
from instructor as T, instructor as S
where T.salary > S.salary and S.dept_name = 'Comp. Sci.'
```

- Keyword **as** is optional and may be omitted  
*instructor as T  $\equiv$  instructor T*

@SKG

2/18/2020 23

## Self Join Example

- Relation *emp-super*

| person | supervisor |
|--------|------------|
| Bob    | Alice      |
| Mary   | Susan      |
| Alice  | David      |
| David  | Mary       |

- Find the supervisor of "Bob"
- Find the supervisor of the supervisor of "Bob"
- Find ALL the supervisors (direct and indirect) of "Bob"

@SKG

2/18/2020 24

## String Operations

- SQL includes a string-matching operator for comparisons on character strings. The operator **like** uses patterns that are described using two special characters:

- percent (%). **The % character matches any substring.**
- underscore (\_). **The \_ character matches any character.**

- Find the names of all instructors whose name includes the substring "dar".

```
select name
from instructor
where name like '%dar%'
```

- Match the string "100%"

```
like '100 \%' escape '\'
```

in that above we use backslash (\) as the escape character.

@SKG

2/18/2020 25

## String Operations (Cont.)

- Patterns are case sensitive.
- Pattern matching examples:
  - 'Intro%' matches any string beginning with "Intro".
  - '%Comp%' matches any string containing "Comp" as a substring.
  - '\_\_\_' matches any string of exactly three characters.
  - '\_\_\_%' matches any string of at least three characters.
- SQL supports a variety of string operations such as
  - concatenation (using "||")
  - converting from upper to lower case (and vice versa)
  - finding string length, extracting substrings, etc.

@SKG

2/18/2020 26

## Ordering the Display of Tuples

- List in alphabetic order the names of all instructors

```
select distinct name
from instructor
order by name
```

- We may specify **desc** for descending order or **asc** for ascending order, for each attribute; ascending order is the default.
  - Example: **order by name desc**
- Can sort on multiple attributes
  - Example: **order by dept\_name, name**

@SKG

2/18/2020 27

## Where Clause Predicates

- SQL includes a **between** comparison operator
- Example: Find the names of all instructors with salary between \$90,000 and \$100,000 (that is,  $\geq \$90,000$  and  $\leq \$100,000$ )
  - ```
select name
from instructor
where salary between 90000 and 100000
```
- Tuple comparison
 - ```
select name, course_id
from instructor, teaches
where (instructor.ID, dept_name) = (teaches.ID, 'Biology');
```

@SKG

2/18/2020 28



## Duplicates

- In relations with duplicates, SQL can define how many copies of tuples appear in the result.
- Multiset** versions of some of the relational algebra operators – given multiset relations  $r_1$  and  $r_2$ :
  - $\sigma_\theta(r_1)$ : If there are  $c_1$  copies of tuple  $t_1$  in  $r_1$ , and  $t_1$  satisfies selections  $\sigma_\theta$ , then there are  $c_1$  copies of  $t_1$  in  $\sigma_\theta(r_1)$ .
  - $\Pi_A(r)$ : For each copy of tuple  $t_i$  in  $r$ , there is a copy of tuple  $\Pi_A(t_i)$  in  $\Pi_A(r)$  where  $\Pi_A(t_i)$  denotes the projection of the single tuple  $t_i$ .
  - $r_1 \times r_2$ : If there are  $c_1$  copies of tuple  $t_1$  in  $r_1$  and  $c_2$  copies of tuple  $t_2$  in  $r_2$ , there are  $c_1 \times c_2$  copies of the tuple  $t_1, t_2$  in  $r_1 \times r_2$ .

@SKG

2/18/2020 29

## Duplicates (Cont.)

- Example: Suppose multiset relations  $r_1(A, B)$  and  $r_2(C)$  are as follows:

$$r_1 = \{(1, a), (2, a)\} \quad r_2 = \{(2), (3), (3)\}$$

- Then  $\Pi_B(r_1)$  would be  $\{(a), (a)\}$ , while  $\Pi_B(r_1) \times r_2$  would be  $\{(a, 2), (a, 2), (a, 3), (a, 3), (a, 3), (a, 3)\}$

- SQL duplicate semantics:

```
select A1, A2, ..., An
from r1, r2, ..., rm
where P
```

is equivalent to the *multiset* version of the expression:

$$\Pi_{A_1, A_2, \dots, A_n}(\sigma_P(r_1 \times r_2 \times \dots \times r_m))$$

@SKG

2/18/2020 30

## Set Operations

- Find courses that ran in Fall 2009 or in Spring 2010  
 (select course\_id from section where sem = 'Fall' and year = 2009)  
 union  
 (select course\_id from section where sem = 'Spring' and year = 2010)
- Find courses that ran in Fall 2009 and in Spring 2010  
 (select course\_id from section where sem = 'Fall' and year = 2009)  
 intersect  
 (select course\_id from section where sem = 'Spring' and year = 2010)
- Find courses that ran in Fall 2009 but not in Spring 2010  
 (select course\_id from section where sem = 'Fall' and year = 2009)  
 except  
 (select course\_id from section where sem = 'Spring' and year = 2010)

@SKG

2/18/2020 31

## Set Operations (Cont.)

- Find the salaries of all instructors that are less than the largest salary.
  - select distinct T.salary  
 from instructor as T, instructor as S  
 where T.salary < S.salary
- Find all the salaries of all instructors
  - select distinct salary  
 from instructor
- Find the largest salary of all instructors.
  - (select "second query")  
 except  
 (select "first query")

@SKG

2/18/2020 32



## Set Operations (Cont.)

- Set operations **union**, **intersect**, and **except**
  - Each of the above operations automatically eliminates duplicates
- To retain all duplicates use the corresponding multiset versions **union all**, **intersect all** and **except all**.
- Suppose a tuple occurs  $m$  times in  $r$  and  $n$  times in  $s$ , then, it occurs:
  - $m + n$  times in  **$r$  union all  $s$**
  - $\min(m, n)$  times in  **$r$  intersect all  $s$**
  - $\max(0, m - n)$  times in  **$r$  except all  $s$**

@SKG

2/18/2020 33

## Null Values

- It is possible for tuples to have a null value, denoted by *null*, for some of their attributes
- *null* signifies an unknown value or that a value does not exist.
- The result of any arithmetic expression involving *null* is *null*
  - Example:  $5 + \text{null}$  returns null
- The predicate **is null** can be used to check for null values.
  - Example: Find all instructors whose salary is null.

```
select name
from instructor
where salary is null
```

@SKG

2/18/2020 34

## Null Values and Three Valued Logic

- Three values – *true*, *false*, *unknown*
- Any comparison with *null* returns *unknown*
  - Example:  $5 < \text{null}$  or  $\text{null} < \text{null}$  or  $\text{null} = \text{null}$
- Three-valued logic using the value *unknown*:
  - OR:  $(\text{unknown or true}) = \text{true}$ ,  
 $(\text{unknown or false}) = \text{unknown}$ ,  
 $(\text{unknown or unknown}) = \text{unknown}$
  - AND:  $(\text{true and unknown}) = \text{unknown}$ ,  
 $(\text{false and unknown}) = \text{false}$ ,  
 $(\text{unknown and unknown}) = \text{unknown}$
  - NOT:  $(\text{not unknown}) = \text{unknown}$
  - “ $P$  is unknown” evaluates to true if predicate  $P$  evaluates to *unknown*
- Result of **where** clause predicate is treated as *false* if it evaluates to *unknown*

@SKG

2/18/2020 35

## Aggregate Functions

- These functions operate on the multiset of values of a column of a relation, and return a value

**avg**: average value  
**min**: minimum value  
**max**: maximum value  
**sum**: sum of values  
**count**: number of values

@SKG

2/18/2020 36

## Aggregate Functions (Cont.)

- Find the average salary of instructors in the Computer Science department
  - `select avg (salary)`  
`from instructor`  
`where dept_name= 'Comp. Sci.';`
- Find the total number of instructors who teach a course in the Spring 2010 semester
  - `select count (distinct ID)`  
`from teaches`  
`where semester = 'Spring' and year = 2010;`
- Find the number of tuples in the *course* relation
  - `select count (*)`  
`from course;`  
@SKG

2/18/2020 37

## Aggregate Functions – Group By

- Find the average salary of instructors in each department
  - `select dept_name, avg (salary) as avg_salary`  
`from instructor`  
`group by dept_name;`

| ID    | name       | dept_name  | salary |
|-------|------------|------------|--------|
| 76766 | Crick      | Biology    | 72000  |
| 45565 | Katz       | Comp. Sci. | 75000  |
| 10101 | Srinivasan | Comp. Sci. | 65000  |
| 83821 | Brandt     | Comp. Sci. | 92000  |
| 98345 | Kim        | Elec. Eng. | 80000  |
| 12121 | Wu         | Finance    | 90000  |
| 76543 | Singh      | Finance    | 80000  |
| 32343 | El Said    | History    | 60000  |
| 58583 | Califieri  | History    | 62000  |
| 15151 | Mozart     | Music      | 40000  |
| 33456 | Gold       | Physics    | 87000  |
| 22222 | Einstein   | Physics    | 95000  |

| dept_name  | avg_salary |
|------------|------------|
| Biology    | 72000      |
| Comp. Sci. | 77333      |
| Elec. Eng. | 80000      |
| Finance    | 85000      |
| History    | 61000      |
| Music      | 40000      |
| Physics    | 91000      |

@SKG

2/18/2020 38

## Aggregation (Cont.)

- Attributes in **select** clause outside of aggregate functions must appear in **group by** list
  - `/* erroneous query */`  
`select dept_name, ID, avg (salary)`  
`from instructor`  
`group by dept_name;`

@SKG

2/18/2020 39

## Aggregate Functions – Having Clause

- Find the names and average salaries of all departments whose average salary is greater than 42000

```
select dept_name, avg (salary)
from instructor
group by dept_name
having avg (salary) > 42000;
```

Note: predicates in the **having** clause are applied after the formation of groups whereas predicates in the **where** clause are applied before forming groups

@SKG

2/18/2020 40

## Null Values and Aggregates

- Total all salaries

```
select sum (salary)
from instructor
```

- Above statement ignores null amounts
- Result is *null* if there is no non-null amount
- All aggregate operations except **count(\*)** ignore tuples with null values on the aggregated attributes
- What if collection has only null values?
  - count returns 0
  - all other aggregates return null

@SKG

2/18/2020 41

## Nested Subqueries

- SQL provides a mechanism for the nesting of subqueries. A **subquery** is a **select-from-where** expression that is nested within another query.
- The nesting can be done in the following SQL query

```
select $A_{s1}, A_{s2}, \dots, A_n$
from $r_{s1}, r_{s2}, \dots, r_m$
where P
```

as follows:

- $A_i$  can be replaced by a subquery that generates a single value.
- $r_i$  can be replaced by any valid subquery
- $P$  can be replaced with an expression of the form:  
 $B <\text{operation}> (\text{subquery})$

Where  $B$  is an attribute and  $<\text{operation}>$  to be defined later.

@SKG

2/18/2020 42

## Subqueries in the Where Clause

@SKG

2/18/2020 43

## Subqueries in the Where Clause

- A common use of subqueries is to perform tests:
  - For set membership
  - For set comparisons
  - For set cardinality.

@SKG

2/18/2020 44

## Set Membership

- Find courses offered in Fall 2009 and in Spring 2010

```
select distinct course_id
from section
where semester = 'Fall' and year= 2009 and
 course_id in (select course_id
 from section
 where semester = 'Spring' and year= 2010);
```

- Find courses offered in Fall 2009 but not in Spring 2010

```
select distinct course_id
from section
where semester = 'Fall' and year= 2009 and
 course_id not in (select course_id
 from section
 where semester = 'Spring' and year= 2010);
```

@SKG

2/18/2020 45

## Set Membership (Cont.)

- Find the total number of (distinct) students who have taken course sections taught by the instructor with ID 10101

```
select count (distinct ID)
from takes
where (course_id, sec_id, semester, year) in
 (select course_id, sec_id, semester, year
 from teaches
 where teaches.ID= 10101);
```

- Note: Above query can be written in a much simpler manner. The formulation above is simply to illustrate SQL features.

@SKG

2/18/2020 46

## Set Comparison – “some” Clause

- Find names of instructors with salary greater than that of some (at least one) instructor in the Biology department.

```
select distinct T.name
from instructor as T, instructor as S
where T.salary > S.salary and S.dept name = 'Biology';
```

- Same query using > some clause

```
select name
from instructor
where salary > some (select salary
 from instructor
 where dept name = 'Biology');
```

@SKG

2/18/2020 47

## Definition of “some” Clause

- $F < \text{comp} > \text{some } r \Leftrightarrow \exists t \in r \text{ such that } (F < \text{comp} > t)$   
Where <comp> can be: <, ≤, >, =, ≠

$(5 < \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline 6 \\ \hline \end{array}) = \text{true}$  (read: 5 < some tuple in the relation)

$(5 < \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{false}$

$(5 = \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{true}$

$(5 \neq \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{true (since } 0 \neq 5)$

$(= \text{some}) \equiv \text{in}$   
However,  $(\neq \text{some}) \neq \text{not in}$

@SKG

2/18/2020 48

## Set Comparison – “all” Clause

- Find the names of all instructors whose salary is greater than the salary of all instructors in the Biology department.

```
select name
from instructor
where salary > all (select salary
 from instructor
 where dept name = 'Biology');
```

@SKG

2/18/2020 49

## Definition of “all” Clause

- $F < \text{comp} > \text{all } r \Leftrightarrow \forall t \in r (F < \text{comp} > t)$

$(5 < \text{all } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline 6 \\ \hline \end{array}) = \text{false}$

$(5 < \text{all } \begin{array}{|c|} \hline 6 \\ \hline 10 \\ \hline \end{array}) = \text{true}$

$(5 = \text{all } \begin{array}{|c|} \hline 4 \\ \hline 5 \\ \hline \end{array}) = \text{false}$

$(5 \neq \text{all } \begin{array}{|c|} \hline 4 \\ \hline 6 \\ \hline \end{array}) = \text{true (since } 5 \neq 4 \text{ and } 5 \neq 6)$

$(\neq \text{all}) \equiv \text{not in}$   
However,  $(= \text{all}) \neq \text{in}$

@SKG

2/18/2020 50

## Test for Empty Relations

- The **exists** construct returns the value **true** if the argument subquery is nonempty.
- exists**  $r \Leftrightarrow r \neq \emptyset$
- not exists**  $r \Leftrightarrow r = \emptyset$

@SKG

2/18/2020 51

## Use of “exists” Clause

- Yet another way of specifying the query “Find all courses taught in both the Fall 2009 semester and in the Spring 2010 semester”

```
select course_id
from section as S
where semester = 'Fall' and year = 2009 and
exists (select *
 from section as T
 where semester = 'Spring' and year = 2010
 and S.course_id = T.course_id);
```

- Correlation name** – variable S in the outer query
- Correlated subquery** – the inner query

@SKG

2/18/2020 52

## Use of “not exists” Clause

- Find all students who have taken all courses offered in the Biology department.

```
select distinct S.ID, S.name
from student as S
where not exists ((select course_id
 from course
 where dept_name = 'Biology')
except
(select T.course_id
 from takes as T
 where S.ID = T.ID));
```

- First nested query lists all courses offered in Biology
- Second nested query lists all courses a particular student took

■ Note that  $X - Y = \emptyset \Leftrightarrow X \subseteq Y$

■ Note: Cannot write this query using = **all** and its variants

@SKG

2/18/2020 53

## Test for Absence of Duplicate Tuples

- The **unique** construct tests whether a subquery has any duplicate tuples in its result.
- The **unique** construct evaluates to “true” if a given subquery contains no duplicates.
- Find all courses that were offered at most once in 2009

```
select T.course_id
from course as T
where unique (select R.course_id
 from section as R
 where T.course_id= R.course_id
 and R.year = 2009);
```

@SKG

2/18/2020 54

## Subqueries in the Form Clause

@SKG

2/18/2020 55

## Subqueries in the Form Clause

- SQL allows a subquery expression to be used in the **from** clause
- Find the average instructors' salaries of those departments where the average salary is greater than \$42,000."

```
select dept_name, avg_salary
from (select dept_name, avg (salary) as avg_salary
 from instructor
 group by dept_name)
where avg_salary > 42000;
```

- Note that we do not need to use the **having** clause
- Another way to write above query

```
select dept_name, avg_salary
from (select dept_name, avg (salary)
 from instructor
 group by dept_name) as dept_avg (dept_name, avg_salary)
where avg_salary > 42000;
```

@SKG

2/18/2020 56

## With Clause

- The **with** clause provides a way of defining a temporary relation whose definition is available only to the query in which the **with** clause occurs.
- Find all departments with the maximum budget

```
with max_budget (value) as
 (select max(budget)
 from department)
select department.name
from department, max_budget
where department.budget = max_budget.value;
```

@SKG

2/18/2020 57

## Complex Queries using With Clause

- Find all departments where the total salary is greater than the average of the total salary at all departments

```
with dept_total (dept_name, value) as
 (select dept_name, sum(salary)
 from instructor
 group by dept_name),
dept_total_avg(value) as
 (select avg(value)
 from dept_total)
select dept_name
from dept_total, dept_total_avg
where dept_total.value > dept_total_avg.value;
```

@SKG

2/18/2020 58

## Scalar Subquery

- Scalar subquery is one which is used where a single value is expected
- List all departments along with the number of instructors in each department

```
select dept_name,
 (select count(*)
 from instructor
 where department.dept_name = instructor.dept_name)
as num_instructors
from department;
```

- Runtime error if subquery returns more than one result tuple

@SKG

2/18/2020 59

## Modification of the Database

- Deletion of tuples from a given relation.
- Insertion of new tuples into a given relation
- Updating of values in some tuples in a given relation

@SKG

2/18/2020 60



## Deletion

- Delete all instructors  
**delete from** *instructor*
- Delete all instructors from the Finance department  
**delete from** *instructor*  
**where** *dept\_name* = 'Finance';
- Delete all tuples in the *instructor* relation for those instructors associated with a department located in the Watson building.

```
delete from instructor
where dept_name in (select dept_name
 from department
 where building = 'Watson');
```

@SKG

2/18/2020 61

## Deletion (Cont.)

- Delete all instructors whose salary is less than the average salary of instructors  
**delete from** *instructor*  
**where** *salary* < (**select** **avg** (*salary*)  
 **from** *instructor*);
- Problem: as we delete tuples from deposit, the average salary changes
- Solution used in SQL:
  1. First, compute **avg** (*salary*) and find all tuples to delete
  2. Next, delete all tuples found above (without recomputing **avg** or retesting the tuples)

@SKG

2/18/2020 62

## Insertion

- Add a new tuple to *course*  
**insert into** *course*  
**values** ('CS-437', 'Database Systems', 'Comp. Sci.', 4);
- or equivalently  
**insert into** *course* (*course\_id*, *title*, *dept\_name*, *credits*)  
**values** ('CS-437', 'Database Systems', 'Comp. Sci.', 4);
- Add a new tuple to *student* with *tot\_creds* set to null  
**insert into** *student*  
**values** ('3003', 'Green', 'Finance', null);

@SKG

2/18/2020 63

## Insertion (Cont.)

- Add all instructors to the *student* relation with *tot\_creds* set to 0  
**insert into** *student*  
**select** *ID*, *name*, *dept\_name*, 0  
**from** *instructor*
- The **select from where** statement is evaluated fully before any of its results are inserted into the relation.  
Otherwise queries like  
**insert into** *table1* **select** \* **from** *table1*  
would cause problem

@SKG

2/18/2020 64

## Updates

- Increase salaries of instructors whose salary is over \$100,000 by 3%, and all others by a 5%

- Write two **update** statements:

```
update instructor
set salary = salary * 1.03
where salary > 100000;
update instructor
set salary = salary * 1.05
where salary <= 100000;
```

- The order is important
  - Can be done better using the **case** statement (next slide)

@SKG

2/18/2020 65

## Case Statement for Conditional Updates

- Same query as before but with case statement

```
update instructor
set salary = case
when salary <= 100000 then salary * 1.05
else salary * 1.03
end
```

@SKG

2/18/2020 66

## Updates with Scalar Subqueries

- Recompute and update tot\_creds value for all students

```
update student S
set tot_cred = (select sum(credits)
from takes, course
where takes.course_id = course.course_id and
S.ID = takes.ID and takes.grade <> 'F' and
takes.grade is not null);
```

- Sets tot\_creds to null for students who have not taken any course
  - Instead of sum(credits), use:

```
case
when sum(credits) is not null then sum(credits)
else 0
end
```

@SKG

2/18/2020 67

## Joined Relations

- **Join operations** take two relations and return as a result another relation.
- A join operation is a Cartesian product which requires that tuples in the two relations match (under some condition). It also specifies the attributes that are present in the result of the join
- The join operations are typically used as subquery expressions in the **from** clause

@SKG

2/18/2020 68

## Types of Join

- Cross Join (Cartesian Product)
- Inner Join
  - Equi-Join
  - Natural Join
- Outer Join
  - Left outer join
  - Right outer join
  - Full outer join
- Self-join

@SKG

2/8/2020 69

## Cross Join

- CROSS JOIN returns the Cartesian product of rows from tables in the join
- Explicit
 

```
select *
 from employee cross join department;
```
- Implicit
 

```
select * from employee, department;
```

@SKG

2/8/2020 70

## Join operations – Example

- Relation *course*

| course_id | title       | dept_name  | credits |
|-----------|-------------|------------|---------|
| BIO-301   | Genetics    | Biology    | 4       |
| CS-190    | Game Design | Comp. Sci. | 4       |
| CS-315    | Robotics    | Comp. Sci. | 3       |

- Relation *prereq*

| course_id | prereq_id |
|-----------|-----------|
| BIO-301   | BIO-101   |
| CS-190    | CS-101    |
| CS-347    | CS-101    |

- Observe that

prereq information is missing for CS-315 and  
course information is missing for CS-437

@SKG

2/8/2020 71

## Inner Join



- *course* **inner join** *prereq*

| course_id | title       | dept_name  | credits | prereq_id | course_id |
|-----------|-------------|------------|---------|-----------|-----------|
| BIO-301   | Genetics    | Biology    | 4       | BIO-101   | BIO-301   |
| CS-190    | Game Design | Comp. Sci. | 4       | CS-101    | CS-190    |

- If specified as **natural**, the 2<sup>nd</sup> course\_id field is skipped

| course_id | title       | dept_name  | credits |
|-----------|-------------|------------|---------|
| BIO-301   | Genetics    | Biology    | 4       |
| CS-190    | Game Design | Comp. Sci. | 4       |
| CS-315    | Robotics    | Comp. Sci. | 3       |

| course_id | prereq_id |
|-----------|-----------|
| BIO-301   | BIO-101   |
| CS-190    | CS-101    |
| CS-347    | CS-101    |

## Outer Join

- An extension of the join operation that avoids loss of information.
- Computes the join and then adds tuples from one relation that does not match tuples in the other relation to the result of the join.
- Uses *null* values.


@SKG

2/18/2020 73

## Left Outer Join

■ *course* **natural left outer join** *prereq*

| <i>course_id</i> | <i>title</i> | <i>dept_name</i> | <i>credits</i> | <i>prereq_id</i> |
|------------------|--------------|------------------|----------------|------------------|
| BIO-301          | Genetics     | Biology          | 4              | BIO-101          |
| CS-190           | Game Design  | Comp. Sci.       | 4              | CS-101           |
| CS-315           | Robotics     | Comp. Sci.       | 3              | <i>null</i>      |



| <i>course_id</i> | <i>title</i> | <i>dept_name</i> | <i>credits</i> | <i>course_id</i> | <i>prereq_id</i> |
|------------------|--------------|------------------|----------------|------------------|------------------|
| BIO-301          | Genetics     | Biology          | 4              | BIO-301          | BIO-101          |
| CS-190           | Game Design  | Comp. Sci.       | 4              | CS-190           | CS-101           |
| CS-315           | Robotics     | Comp. Sci.       | 3              | CS-347           | CS-101           |


@SKG

2/18/2020 74

## Right Outer Join

■ *course* **natural right outer join** *prereq*

| <i>course_id</i> | <i>title</i> | <i>dept_name</i> | <i>credits</i> | <i>prereq_id</i> |
|------------------|--------------|------------------|----------------|------------------|
| BIO-301          | Genetics     | Biology          | 4              | BIO-101          |
| CS-190           | Game Design  | Comp. Sci.       | 4              | CS-101           |
| CS-347           | <i>null</i>  | <i>null</i>      | <i>null</i>    | CS-101           |



| <i>course_id</i> | <i>title</i> | <i>dept_name</i> | <i>credits</i> | <i>course_id</i> | <i>prereq_id</i> |
|------------------|--------------|------------------|----------------|------------------|------------------|
| BIO-301          | Genetics     | Biology          | 4              | BIO-301          | BIO-101          |
| CS-190           | Game Design  | Comp. Sci.       | 4              | CS-190           | CS-101           |
| CS-315           | Robotics     | Comp. Sci.       | 3              | CS-347           | CS-101           |

@SKG

2/18/2020 75

## Joined Relations

- **Join operations** take two relations and return as a result another relation.
- These additional operations are typically used as subquery expressions in the **from** clause
- **Join condition** – defines which tuples in the two relations match, and what attributes are present in the result of the join.
- **Join type** – defines how tuples in each relation that do not match any tuple in the other relation (based on the join condition) are treated.


| <i>Join types</i> | <i>Join Conditions</i>                  |
|-------------------|-----------------------------------------|
| inner join        | <b>natural</b>                          |
| left outer join   | <b>on</b> <predicate>                   |
| right outer join  | <b>using</b> ( $A_1, A_1, \dots, A_n$ ) |
| full outer join   |                                         |

76

## Full Outer Join

- **course natural full outer join prereq**

| course_id | title       | dept_name  | credits | prereq_id |
|-----------|-------------|------------|---------|-----------|
| BIO-301   | Genetics    | Biology    | 4       | BIO-101   |
| CS-190    | Game Design | Comp. Sci. | 4       | CS-101    |
| CS-315    | Robotics    | Comp. Sci. | 3       | null      |
| CS-347    | null        | null       | null    | CS-101    |



| course_id | title       | dept_name  | credits |
|-----------|-------------|------------|---------|
| BIO-301   | Genetics    | Biology    | 4       |
| CS-190    | Game Design | Comp. Sci. | 4       |
| CS-315    | Robotics    | Comp. Sci. | 3       |

| course_id | prereq_id |
|-----------|-----------|
| BIO-301   | BIO-101   |
| CS-190    | CS-101    |
| CS-347    | CS-101    |

@SKG

2/18/2020 77

## Joined Relations – Examples

- **course inner join prereq on**  
*course.course\_id = prereq.course\_id*

| course_id | title       | dept_name  | credits | prereq_id | course_id |
|-----------|-------------|------------|---------|-----------|-----------|
| BIO-301   | Genetics    | Biology    | 4       | BIO-101   | BIO-301   |
| CS-190    | Game Design | Comp. Sci. | 4       | CS-101    | CS-190    |

- What is the difference between the above, and a natural join?

- **course left outer join prereq on**  
*course.course\_id = prereq.course\_id*

| course_id | title       | dept_name  | credits | prereq_id | course_id |
|-----------|-------------|------------|---------|-----------|-----------|
| BIO-301   | Genetics    | Biology    | 4       | BIO-101   | BIO-301   |
| CS-190    | Game Design | Comp. Sci. | 4       | CS-101    | CS-190    |
| CS-315    | Robotics    | Comp. Sci. | 3       | null      | null      |

@SKG

2/18/2020 78

## Joined Relations – Examples

- **course natural right outer join prereq**

| course_id | title       | dept_name  | credits | prereq_id |
|-----------|-------------|------------|---------|-----------|
| BIO-301   | Genetics    | Biology    | 4       | BIO-101   |
| CS-190    | Game Design | Comp. Sci. | 4       | CS-101    |
| CS-347    | null        | null       | null    | CS-101    |

- **course full outer join prereq using (course\_id)**

| course_id | title       | dept_name  | credits | prereq_id |
|-----------|-------------|------------|---------|-----------|
| BIO-301   | Genetics    | Biology    | 4       | BIO-101   |
| CS-190    | Game Design | Comp. Sci. | 4       | CS-101    |
| CS-315    | Robotics    | Comp. Sci. | 3       | null      |
| CS-347    | null        | null       | null    | CS-101    |

@SKG

2/18/2020 79

## Views

- In some cases, it is not desirable for all users to see the entire logical model (that is, all the actual relations stored in the database.)
- Consider a person who needs to know an instructors name and department, but not the salary. This person should see a relation described, in SQL, by

```
select ID, name, dept_name
from instructor
```

- A **view** provides a mechanism to hide certain data from the view of certain users.
- Any relation that is not of the conceptual model but is made visible to a user as a "virtual relation" is called a **view**.

@SKG

2/18/2020 80

## View Definition

- A view is defined using the **create view** statement which has the form

**create view v as** < query expression >

where <query expression> is any legal SQL expression. The view name is represented by v.

- Once a view is defined, the view name can be used to refer to the virtual relation that the view generates.
- View definition is not the same as creating a new relation by evaluating the query expression
  - Rather, a view definition causes the saving of an expression; the expression is substituted into queries using the view.

@SKG

2/18/2020 81

## Example Views

- A view of instructors without their salary  
**create view faculty as**  
**select ID, name, dept\_name**  
**from instructor**
- Find all instructors in the Biology department  
**select name**  
**from faculty**  
**where dept\_name = 'Biology'**
- Create a view of department salary totals  
**create view departments\_total\_salary(dept\_name, total\_salary) as**  
**select dept\_name, sum(salary)**  
**from instructor**  
**group by dept\_name;**

@SKG

2/18/2020 82

## Views Defined Using Other Views

- create view physics\_fall\_2009 as**  
**select course.course\_id, sec\_id, building, room\_number**  
**from course, section**  
**where course.course\_id = section.course\_id**  
**and course.dept\_name = 'Physics'**  
**and section.semester = 'Fall'**  
**and section.year = '2009';**
- create view physics\_fall\_2009\_watson as**  
**select course\_id, room\_number**  
**from physics\_fall\_2009**  
**where building = 'Watson';**

@SKG

2/18/2020 83

## View Expansion

- Expand use of a view in a query/another view  
**create view physics\_fall\_2009\_watson as**  
**(select course\_id, room\_number**  
**from (select course.course\_id, building, room\_number**  
**from course, section**  
**where course.course\_id = section.course\_id**  
**and course.dept\_name = 'Physics'**  
**and section.semester = 'Fall'**  
**and section.year = '2009')**  
**where building = 'Watson';**

@SKG

2/18/2020 84

## Views Defined Using Other Views

- One view may be used in the expression defining another view
- A view relation  $v_1$  is said to *depend directly* on a view relation  $v_2$  if  $v_2$  is used in the expression defining  $v_1$
- A view relation  $v_1$  is said to *depend on* view relation  $v_2$  if either  $v_1$  depends directly to  $v_2$  or there is a path of dependencies from  $v_1$  to  $v_2$
- A view relation  $v$  is said to be *recursive* if it depends on itself.

@SKG

2/18/2020 85

## View Expansion

- A way to define the meaning of views defined in terms of other views.
- Let view  $v_1$  be defined by an expression  $e_1$  that may itself contain uses of view relations.
- View expansion of an expression repeats the following replacement step:
  - repeat**
    - Find any view relation  $v_i$  in  $e_1$
    - Replace the view relation  $v_i$  by the expression defining  $v_i$
  - until** no more view relations are present in  $e_1$
- As long as the view definitions are not recursive, this loop will terminate

@SKG

2/18/2020 86

## Recursive View

- In SQL, recursive queries are typically built using these components:
  - A non-recursive seed statement
  - A recursive statement
  - A connection operator
  - The only valid set connection operator in a recursive view definition is UNION ALL
  - A terminal condition to prevent infinite recursion

@SKG

2/18/2020 87

## Recursive View- Example

- In the context of a relation *flights*:

| create table flights ( |                |          |             |                   |
|------------------------|----------------|----------|-------------|-------------------|
| source                 | varchar(40),   | source   | destination | carrier           |
| destination            | varchar(40),   | Paris    | Detroit     | KLM               |
| carrier                | varchar(40),   | Paris    | New York    | KLM               |
| cost                   | decimal(5,0)); | Paris    | Boston      | American Airlines |
|                        |                | New York | Chicago     | American Airlines |
|                        |                | Boston   | Chicago     | American Airlines |
|                        |                | Detroit  | San Jose    | American Airlines |
|                        |                | Chicago  | San Jose    | American Airlines |

- Find all the destinations that can be reached from 'Paris'

@SKG

2/18/2020 88



## Recursive View- Example

```
create recursive view reachable_from (source,destination,depth) as (
 select root.source, root.destination, 0 as depth
 from flights as root
 where root.source = 'Paris'
 union all
 select in1.source, out1.destination, in1.depth + 1
 from reachable_from as in1, flights as out1
 where in1.destination = out1.source and
 in1.depth <= 100);
```

- Get the result by simple selection on the view:

```
select distinct source, destination
from reachable_from;
```

This example view, *reachable\_from*, is called the *transitive closure* of the *flights* relation

| source | destination |
|--------|-------------|
| Paris  | Detroit     |
| Paris  | New York    |
| Paris  | Boston      |
| Paris  | Chicago     |
| Paris  | San Jose    |

\*A non-recursive seed statement  
 \*A recursive statement  
 \*A connection operator  
 \*A terminal condition to prevent infinite recursion

| source   | destination | carrier           | cost |
|----------|-------------|-------------------|------|
| Paris    | Detroit     | KLM               | 7    |
| Paris    | New York    | KLM               | 6    |
| Paris    | Boston      | American Airlines | 8    |
| New York | Chicago     | American Airlines | 2    |
| Boston   | Chicago     | American Airlines | 6    |
| Detroit  | San Jose    | American Airlines | 4    |
| Chicago  | San Jose    | American Airlines | 2    |

@SKG

2/18/2020 89

## The power of Recursion

- Recursive views make it possible to write queries, such as transitive closure queries, that cannot be written without recursion or iteration
  - Intuition: Without recursion, a non-recursive non-iterative program can perform only a fixed number of joins of *flights* with itself
    - This can give only a fixed number of levels of reachable destinations
    - Given a fixed non-recursive query, we can construct a database with a greater number of levels of reachable destinations on which the query will not work

@SKG

2/18/2020 90

## The power of Recursion

- Computing transitive closure using iteration, adding successive tuples to *reachable\_from*
  - The next slide shows a *flights* relation
  - Each step of the iterative process constructs an extended version of *reachable\_from* from its recursive definition
  - The final result is called the *fixed point* of the recursive view definition.
- Recursive views are required to be **monotonic**. That is, if we add tuples to *flights* the view *reachable\_from* contains all of the tuples it contained before, plus possibly more

@SKG

2/18/2020 91

## Fixed point Computation: Example

| source   | destination | carrier           | cost |
|----------|-------------|-------------------|------|
| Paris    | Detroit     | KLM               | 7    |
| Paris    | New York    | KLM               | 6    |
| Paris    | Boston      | American Airlines | 8    |
| New York | Chicago     | American Airlines | 2    |
| Boston   | Chicago     | American Airlines | 6    |
| Detroit  | San Jose    | American Airlines | 4    |
| Chicago  | San Jose    | American Airlines | 2    |

| Iteration # | Tuples in Closure                            |
|-------------|----------------------------------------------|
| 0           | Detroit, New York, Boston                    |
| 1           | Detroit, New York, Boston, San Jose, Chicago |
| 2           | Detroit, New York, Boston, San Jose, Chicago |

@SKG

2/18/2020 92

## Update of a View

- Add a new tuple to *faculty* view which we defined earlier  
**insert into faculty values** ('30765', 'Green', 'Music');  
This insertion must be represented by the insertion of the tuple  
('30765', 'Green', 'Music', null)  
into the *instructor* relation

@SKG

2/18/2020 93

## Some Updates cannot be Translated Uniquely

- **create view** *instructor\_info* as  
**select** *ID, name, building*  
**from** *instructor, department*  
**where** *instructor.dept\_name= department.dept\_name;*
- **insert into** *instructor\_info* **values** ('69987', 'White', 'Taylor');
  - ✓ which department, if multiple departments in Taylor?
  - ✓ what if no department is in Taylor?
- Most SQL implementations allow updates only on simple views
  - ✓ The **from** clause has only one database relation.
  - ✓ The **select** clause contains only attribute names of the relation, and does not have any expressions, aggregates, or **distinct** specification.
  - ✓ Any attribute not listed in the **select** clause can be set to null

@SKG

2/18/2020 94

## And Some Not at All

- **create view** *history\_instructors* as  
**select** \*  
**from** *instructor*  
**where** *dept\_name= 'History';*
- What happens if we insert ('25566', 'Brown', 'Biology', 100000) into *history\_instructors*?

@SKG

2/18/2020 95

## Materialized Views

- **Materializing a view:** create a physical table containing all the tuples in the result of the query defining the view
- If relations used in the query are updated, the materialized view result becomes out of date
  - Need to **maintain** the view, by updating the view whenever the underlying relations are updated.

@SKG

2/18/2020 96

## Transactions

- Unit of work
- Atomic transaction
  - either fully executed or rolled back as if it never occurred
- Isolation from concurrent transactions
- Transactions begin implicitly
  - Ended by **commit work** or **rollback work**
- But default on most databases: each SQL statement commits automatically
  - Can turn off auto commit for a session (e.g. using API)
  - In SQL:1999, can use: **begin atomic .... end**
    - ✓ Not supported on most databases

@SKG

2/18/2020 97

## Integrity Constraints

- Integrity constraints guard against accidental damage to the database, by ensuring that authorized changes to the database do not result in a loss of data consistency.
  - ✓ A checking account must have a balance greater than \$10,000.00
  - ✓ A salary of a bank employee must be at least \$4.00 an hour
  - ✓ A customer must have a (non-null) phone number

@SKG

2/18/2020 98

## Integrity Constraints on a Single Relation

- not null
- primary key
- unique
- check (P), where P is a predicate

@SKG

2/18/2020 99

## Not Null and Unique Constraints

- not null
  - Declare *name* and *budget* to be **not null**  
*name* **varchar(20) not null**  
*budget* **numeric(12,2) not null**
- unique (  $A_1, A_2, \dots, A_m$  )
  - The unique specification states that the attributes  $A_1, A_2, \dots, A_m$  form a candidate key.
  - Candidate keys are permitted to be null (in contrast to primary keys).

@SKG

2/18/2020 100

## The check clause

- **check (P)**  
where P is a predicate  
e.g.: ensure that semester is one of fall, winter, spring or summer:

```
create table section (
 course_id varchar (8),
 sec_id varchar (8),
 semester varchar (6),
 year numeric (4,0),
 building varchar (15),
 room_number varchar (7),
 time_slot_id varchar (4),
 primary key (course_id, sec_id, semester, year),
 check (semester in ('Fall', 'Winter', 'Spring', 'Summer'))
);
```

@SKG

2/18/2020 101

## Referential Integrity

- Ensures that a value that appears in one relation for a given set of attributes also appears for a certain set of attributes in another relation.
  - e.g.: If "Biology" is a department name appearing in one of the tuples in the *instructor* relation, then there exists a tuple in the *department* relation for "Biology".
- Let A be a set of attributes. Let R and S be two relations that contain attributes A and where A is the primary key of S. A is said to be a **foreign key** of R if for any values of A appearing in R these values also appear in S.

@SKG

2/18/2020 102

## Cascading Actions in Referential Integrity

- **create table course (**  
  **course\_id char(5) primary key,**  
  **title varchar(20),**  
  **dept\_name varchar(20) references department**  
  **)**
- **create table course (**  
  ...  
  **dept\_name varchar(20),**  
  **foreign key (dept\_name) references department**  
    **on delete cascade**  
    **on update cascade,**  
  ...  
  **)**
- alternative actions to cascade: **set null, set default**

@SKG

2/18/2020 103

## Integrity Constraint Violation During Transactions

- e.g.  
**create table person (**  
  **ID char(10),**  
  **name char(40),**  
  **mother char(10),**  
  **father char(10),**  
  **primary key ID,**  
  **foreign key father references person,**  
  **foreign key mother references person)**
- How to insert a tuple without causing constraint violation?
  - insert father and mother of a person before inserting person
  - OR, set father and mother to null initially, update after inserting all persons (not possible if father and mother attributes declared to be **not null**)

@SKG

2/18/2020 104

## Complex Check Clauses

- **check** (*time\_slot\_id* in (**select** *time\_slot\_id* **from** *time\_slot*))
  - why not use a foreign key here?
- Every section has at least one instructor teaching the section.
  - how to write this?
- Unfortunately: subquery in check clause not supported by pretty much any database
  - Alternative: triggers (later)
- **create assertion** <assertion-name> **check** <predicate>;
  - Also not supported by anyone

@SKG

2/18/2020 105

## Built-in Data Types in SQL

- **date**: Dates, containing a (4 digit) year, month and date
  - Example: **date** '2005-7-27'
- **time**: Time of day, in hours, minutes and seconds.
  - Example: **time** '09:00:30'    **time** '09:00:30.75'
- **timestamp**: date plus time of day
  - Example: **timestamp** '2005-7-27 09:00:30.75'
- **interval**: period of time
  - Example: **interval** '1' day
  - Subtracting a date/time/timestamp value from another gives an interval value
  - Interval values can be added to date/time/timestamp values

@SKG

2/18/2020 106

## Index Creation

- **create table** *student*  
(*ID* **varchar** (5),  
*name* **varchar** (20) **not null**,  
*dept\_name* **varchar** (20),  
*tot\_cred* **numeric** (3,0) **default** 0,  
**primary key** (*ID*))
  - **create index** *studentID\_index* **on** *student*(*ID*)
  - Indices are data structures used to speed up access to records with specified values for index attributes
    - e.g. **select** \*  
      **from** *student*  
      **where** *ID* = '12345'
- can be executed by using the index to find the required record, without looking at all records of *student*

107

## User-Defined Types

- **create type** construct in SQL creates user-defined type

**create type** *Dollars* **as** **numeric** (12,2) **final**

- **create table** *department*  
(*dept\_name* **varchar** (20),  
*building* **varchar** (15),  
*budget* *Dollars*);

@SKG

2/18/2020 108

## Domains

- **create domain** construct in SQL-92 creates user-defined domain types

```
create domain person_name char(20) not null
```

- Types and domains are similar. Domains can have constraints, such as **not null**, specified on them.
- **create domain** *degree\_level* varchar(10)  
**constraint** *degree\_level\_test*  
**check** (value in ('Bachelors', 'Masters', 'Doctorate'));

@SKG

2/18/2020 109

## Large-Object Types

- Large objects (photos, videos, CAD files, etc.) are stored as a *large object*:
  - **blob**: binary large object -- object is a large collection of uninterpreted binary data (whose interpretation is left to an application outside of the database system)
  - **clob**: character large object -- object is a large collection of character data
- When a query returns a large object, a pointer is returned rather than the large object itself.

@SKG

2/18/2020 110

## Authorization

Forms of authorization on parts of the database:

- **Read** - allows reading, but not modification of data.
- **Insert** - allows insertion of new data, but not modification of existing data.
- **Update** - allows modification, but not deletion of data.
- **Delete** - allows deletion of data.

Forms of authorization to modify the database schema

- **Index** - allows creation and deletion of indices.
- **Resources** - allows creation of new relations.
- **Alteration** - allows addition or deletion of attributes in a relation.
- **Drop** - allows deletion of relations.

@SKG

2/18/2020 111

## Authorization Specification in SQL

- The **grant** statement is used to confer authorization

```
grant <privilege list>
on <relation name or view name> to <user list>
```
- <user list> is:
  - a user-id
  - **public**, which allows all valid users the privilege granted
  - A role (more on this later)
- Granting a privilege on a view does not imply granting any privileges on the underlying relations.
- The grantor of the privilege must already hold the privilege on the specified item (or be the database administrator).

@SKG

2/18/2020 112

## Privileges in SQL

- **select**: allows read access to relation, or the ability to query using the view
  - Example: grant users  $U_1$ ,  $U_2$ , and  $U_3$  **select** authorization on the *instructor* relation:  
**grant select on instructor to  $U_1$ ,  $U_2$ ,  $U_3$**
- **insert**: the ability to insert tuples
- **update**: the ability to update using the SQL update statement
- **delete**: the ability to delete tuples.
- **all privileges**: used as a short form for all the allowable privileges

@SKG

2/18/2020 113

## Revoking Authorization in SQL

- The **revoke** statement is used to revoke authorization.  
**revoke** <privilege list>  
**on** <relation name or view name> **from** <user list>
- Example:  
**revoke select on branch from  $U_1$ ,  $U_2$ ,  $U_3$**
- <privilege-list> may be **all** to revoke all privileges the revokee may hold.
- If <revokee-list> includes **public**, all users lose the privilege except those granted it explicitly.
- If the same privilege was granted twice to the same user by different grantees, the user may retain the privilege after the revocation.
- All privileges that depend on the privilege being revoked are also revoked.

@SKG 2/18/2020 114

## Roles

- **create role** instructor;
- **grant instructor to Amit**;
- Privileges can be granted to roles:
  - **grant select on takes to instructor**;
- Roles can be granted to users, as well as to other roles
  - **create role teaching\_assistant**
  - **grant teaching\_assistant to instructor**;
    - Instructor inherits all privileges of *teaching\_assistant*
- Chain of roles
  - **create role dean**;
  - **grant instructor to dean**;
  - **grant dean to Satoshi**;

2/18/2020 115

## Authorization on Views

- **create view geo\_instructor as**  
**(select \***  
**from instructor**  
**where dept\_name = 'Geology');**
- **grant select on geo\_instructor to geo\_staff**
- Suppose that a *geo\_staff* member issues
  - **select \***  
**from geo\_instructor**;
- What if
  - *geo\_staff* does not have permissions on *instructor*?
  - creator of view did not have some permissions on *instructor*?

@SKG

2/18/2020 116



## Other Authorization Features

- **references** privilege to create foreign key
  - ✓ **grant reference** (*dept\_name*) **on** *department* **to** Mariano;
  - ✓ why is this required?
- transfer of privileges
  - ✓ **grant select on** *department* **to** Amit **with grant option**;
  - ✓ **revoke select on** *department* **from** Amit, Satoshi **cascade**;
  - ✓ **revoke select on** *department* **from** Amit, Satoshi **restrict**;
- Etc. read Section 4.6 for more details we have omitted here.

@SKG

2/28/2020 117