

# Transaction Processing: Theory and Concept

Swarup Kr Ghosh  
swarupg1@gmail.com

# Introduction

---

- ❖ Transaction Processing systems are **system with large database and hundred of users** that are executing database transaction
- ❖ Require **high availability** and **fast response time** for hundreds of concurrent users

# Examples

---

- ❖ Airline/ Railway Reservation Systems
- ❖ Banking Systems (ATM, ...)
- ❖ Trading and Brokerage Systems
- ❖ Hotel / Hospital Systems
- ❖ Standard Commercial Systems


# Single User Versus Multi Users System

---


- ❖ **Single-User System:** At most one user at a time can use the system.
- ❖ **Multiuser System:** Many users can access the system concurrently.
- ❖ **Concurrency**
  - **Interleaved processing:** concurrent execution of processes is interleaved in a single CPU
  - **Parallel processing:** processes are concurrently executed in multiple CPUs

# Single User Versus Multi User System

## ❖ Interleaved Process



## ❖ Parallel Process



# Transaction, Read and Write Operations

---

- ❖ A transaction is an **executing program that forms a *logical unit* of database processing**
- ❖ Transaction includes one or more database access operation
  - ❑ Insertion
  - ❑ Deletion
  - ❑ Modification
  - ❑ Retrieval
- ❖ Database operation that form a transaction can be embedded within an application program or can be specified interactively via high level query language like SQL

# Transaction, Read and Write Operations

---

- ❖ **Transaction boundaries:** Begin and End transaction
- ❖ An **application program** may contain several transactions separated by the Begin and End transaction boundaries
- ❖ **Read-only Transaction**
- ❖ A database is specified as a collection of **named data items**
- ❖ The size of a data item is called its **granularity**

# Read and Write Operations

---

- ❖ Basic database access operations include
  - **read\_item (x)** : Reads a database item named x into a program variable
  - **write\_item (x)** : Writes the value of program variable x into the database item named x
- ❖ Basic unit of data transfer from the disk to the computer main memory is one block
- ❖ In general, a data item (what is read or written) will be the field of some record in the database, although it may be a larger unit such as a record or even a whole block



# Read and Write Operations

---

❖ **read\_item (x) includes the following steps:**

1. Find the **address** of the disk block that contains item x
2. **Copy** that disk block into a **buffer in main memory** (if that disk block is not already in some main memory buffer)
3. **Copy** item x from the **buffer to the program variable** named x

# Read and Write Operations

---

❖ **write\_item (x) includes the following steps:**

1. Find the **address** of the disk block that contains item x
2. Copy that **disk block** into a buffer in **main memory** (if that disk block is not already in some main memory buffer)
3. Copy item x from the **program variable** named x into its correct location in the **buffer**
4. Store the **updated block** from the buffer back **to disk** (either immediately or at some later point in time)

# Two Sample Transactions

---

(a)  $T_1$

---

read\_item ( $X$ );  
 $X := X - N$ ;  
write\_item ( $X$ );  
read\_item ( $Y$ );  
 $Y := Y + N$ ;  
write\_item ( $Y$ );

(a) Transaction T1

(b)  $T_2$

---

read\_item ( $X$ );  
 $X := X + M$ ;  
write\_item ( $X$ );

(b) Transaction T2

# Example of Fund Transfer

---

❖ Transaction to transfer \$50 from account A to account B:

1.        **read**(A)
2.         $A := A - 50$
3.        **write**(A)
4.        **read**(B)
5.         $B := B + 50$
6.        **write**(B)

- ❖ if the **transaction fails after step 3 and before step 6**, the system should ensure that its updates are not reflected in the database, else an inconsistency will result
- ❖ the sum of A and B should be unchanged by the execution of the transaction

# Example of Fund Transfer (Contd.)

---

- ❖ If between steps 3 and 6, another transaction is allowed to access the partially updated database, it will see an inconsistent database (the sum  $A + B$  will be less than it should be)
  - One solution, running transactions **serially**, that is one after the other
  - However, executing multiple transactions concurrently has significant benefits, as we will see later
- ❖ Once the user has been notified that the transaction has completed (i.e., the transfer of the \$50 has taken place), the updates to the database by the **transaction must persist despite failures**

# Why Concurrency Control is needed

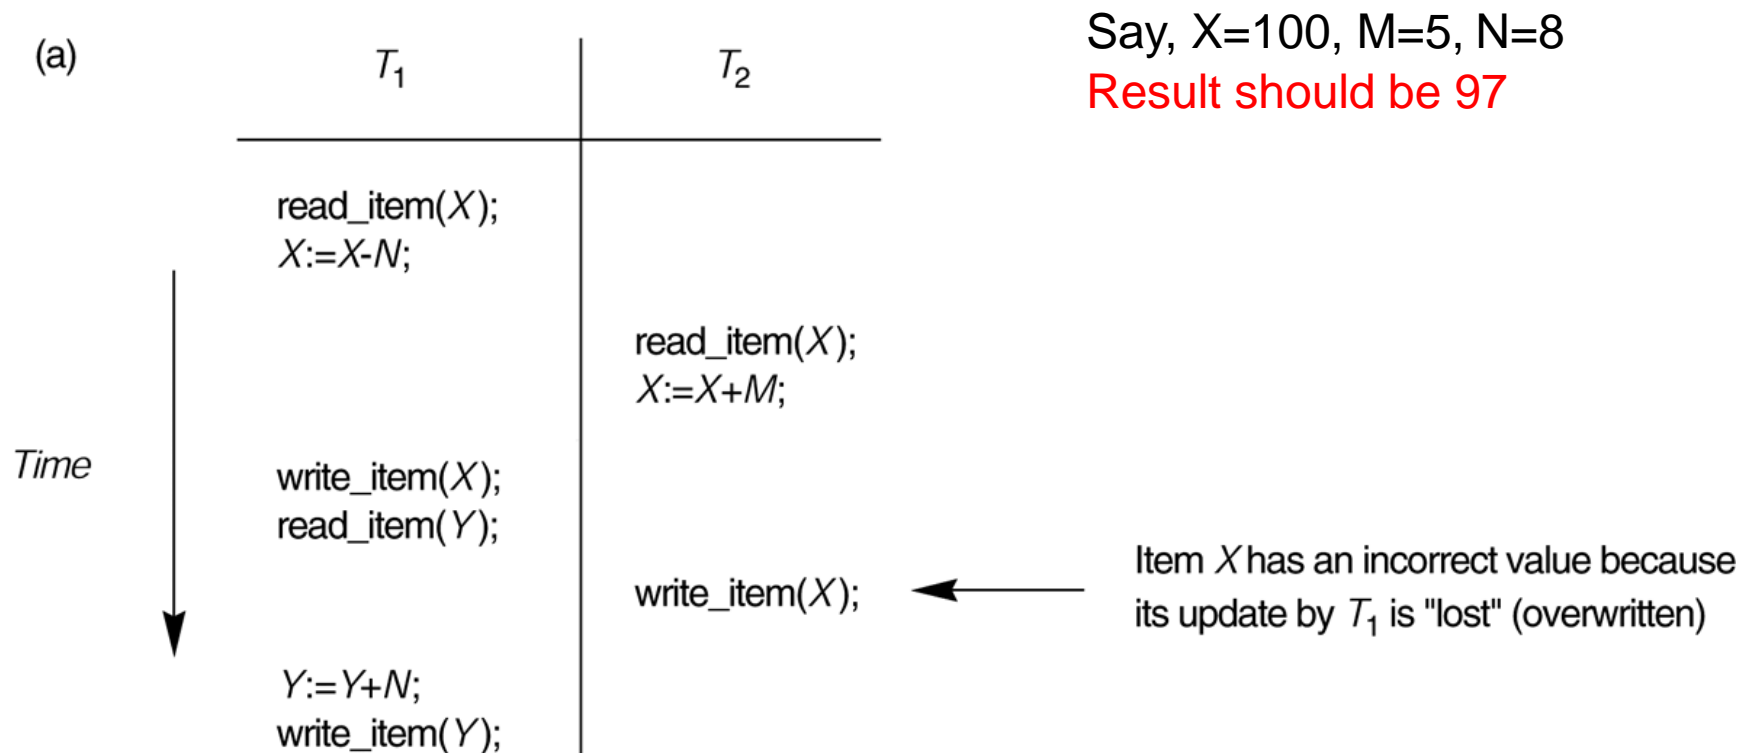
---

## ❖ **Lost Update Problem**

This occurs when two transactions that access the same database items have their operations interleaved in a way that makes the value of some database item incorrect

# Why Concurrency Control is needed

## ❖ Lost Update Problem



# Why Concurrency Control is needed

---

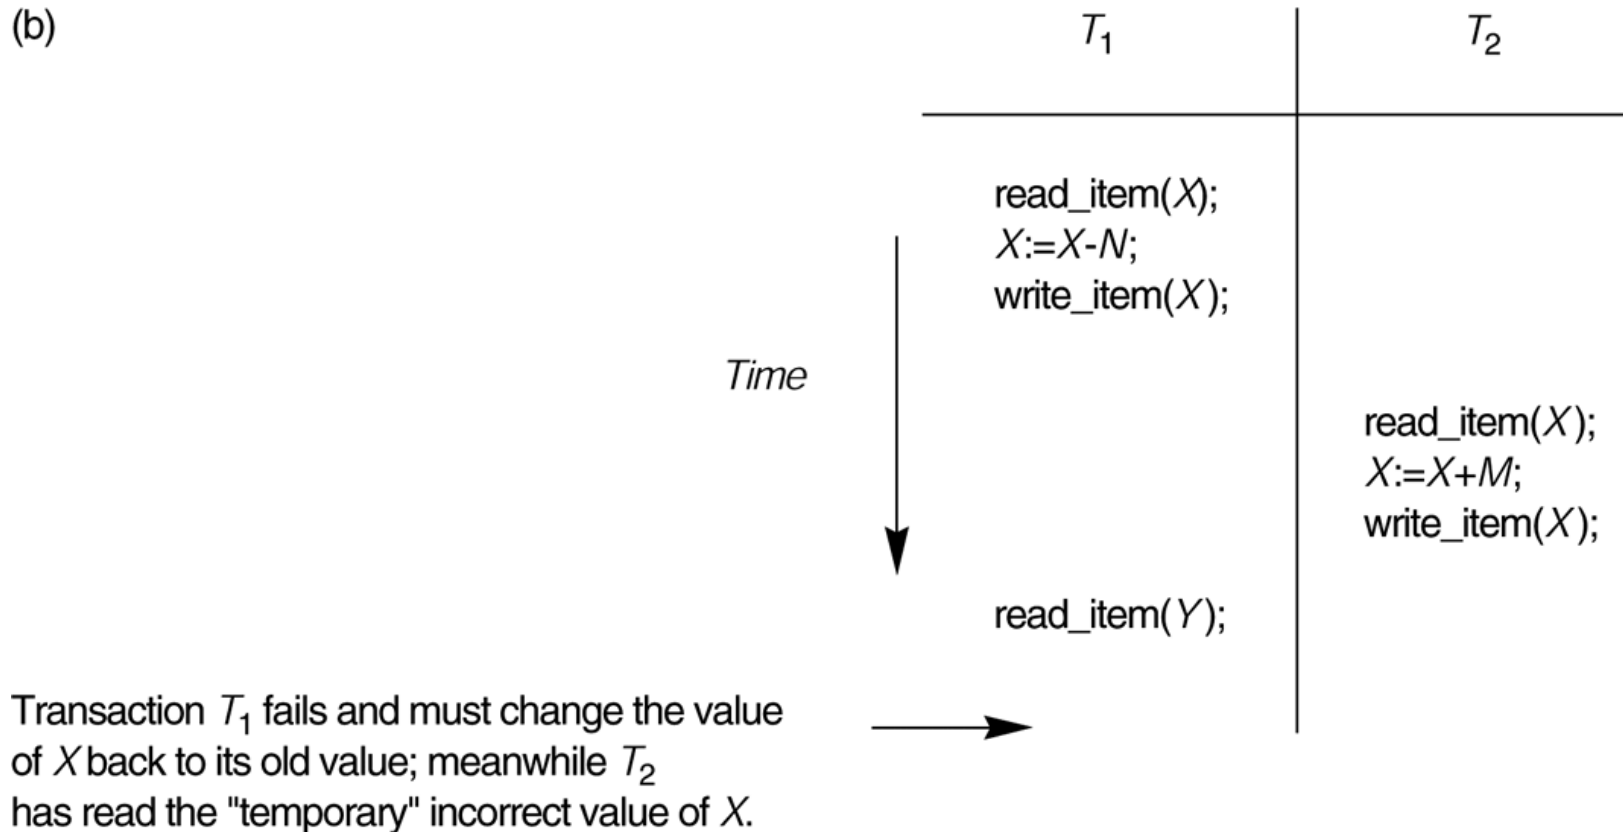
## ❖ **Temporary Update (Dirty Read) Problem**

- This occurs when one transaction updates a database item and then the transaction fails for some reason
- The updated item is accessed by another transaction before it is changed back to its original value



# Why Concurrency Control is needed

## ❖ Temporary Update (Dirty Read) Problem



# Why Concurrency Control is needed

---

## ❖ **Incorrect Summary Problem**

If one transaction is calculating an aggregate summary function on a number of records while other transactions are updating some of these records, the aggregate function may calculate some values before they are updated and others after they are updated

# Why Concurrency Control is needed

## ❖ Incorrect Summary Problem

(c)	$T_1$	$T_3$
		$sum:=0;$ $read\_item(A);$ $sum:=sum+A;$  $\vdots$
	$read\_item(X);$ $X:=X-N;$ $write\_item(X);$	$read\_item(X);$ $sum:=sum+X;$ $read\_item(Y);$ $sum:=sum+Y;$
	$read\_item(Y);$ $Y:=Y+N;$ $write\_item(Y);$	

←  $T_3$  reads  $X$  after  $N$  is subtracted and reads  $Y$  before  $N$  is added; a wrong summary is the result (off by  $N$ ).

# Why Concurrency Control is needed

---

## ❖ **Non-repeatable read**

Transaction T reads an item twice and the item is changed by another transaction T' between the two reads → T receives two different values for its two reads of the same item

Eg. Airlines reservation system

# Transaction Support in SQL

---

- ❖ **Dirty Read**
- ❖ **Non-repeatable read**
- ❖ **Phantom Read**

# Significance of Recovery

---

- ❖ What causes a Transaction to fail
  - ❖ A Computer Failure (system crash)
  - ❖ A Transaction or System Error
  - ❖ Local Errors or Exception Conditions
  - ❖ Concurrency Control Enforcement
  - ❖ Disk Failure
  - ❖ Physical Problems and Catastrophes

# ACID Properties

---

- ❖ **A**tomicity
- ❖ **C**onsistency
- ❖ **I**solation
- ❖ **D**urability

# ACID Properties

---

A **transaction** is a unit of program execution that accesses and possibly updates various data items. To preserve the integrity of data the database system must ensure:

**Atomicity:** Either all operations of the transaction are properly reflected in the database or none are.

**Consistency:** Execution of a transaction in isolation preserves the consistency of the database.

**Isolation:** Although multiple transactions may execute concurrently, each transaction must be unaware of other concurrently executing transactions. Intermediate transaction results must be hidden from other concurrently executed transactions.

That is, for every pair of transactions  $T_i$  and  $T_j$ , it appears to  $T_i$  that either  $T_j$  finished execution before  $T_i$  started, or  $T_j$  started execution after  $T_i$  finished.

**Durability:** After a transaction completes successfully, the changes it has made to the database persist, even if there are system failures.



# Example of Fund Transfer

---

❖ Transaction to transfer \$50 from account A to account B:

1.        **read**(A)
2.         $A := A - 50$
3.        **write**(A)
4.        **read**(B)
5.         $B := B + 50$
6.        **write**(B)

- ❖ if the **transaction fails after step 3 and before step 6**, the system should ensure that its updates are not reflected in the database, else an inconsistency will result
- ❖ the sum of A and B should be unchanged by the execution of the transaction

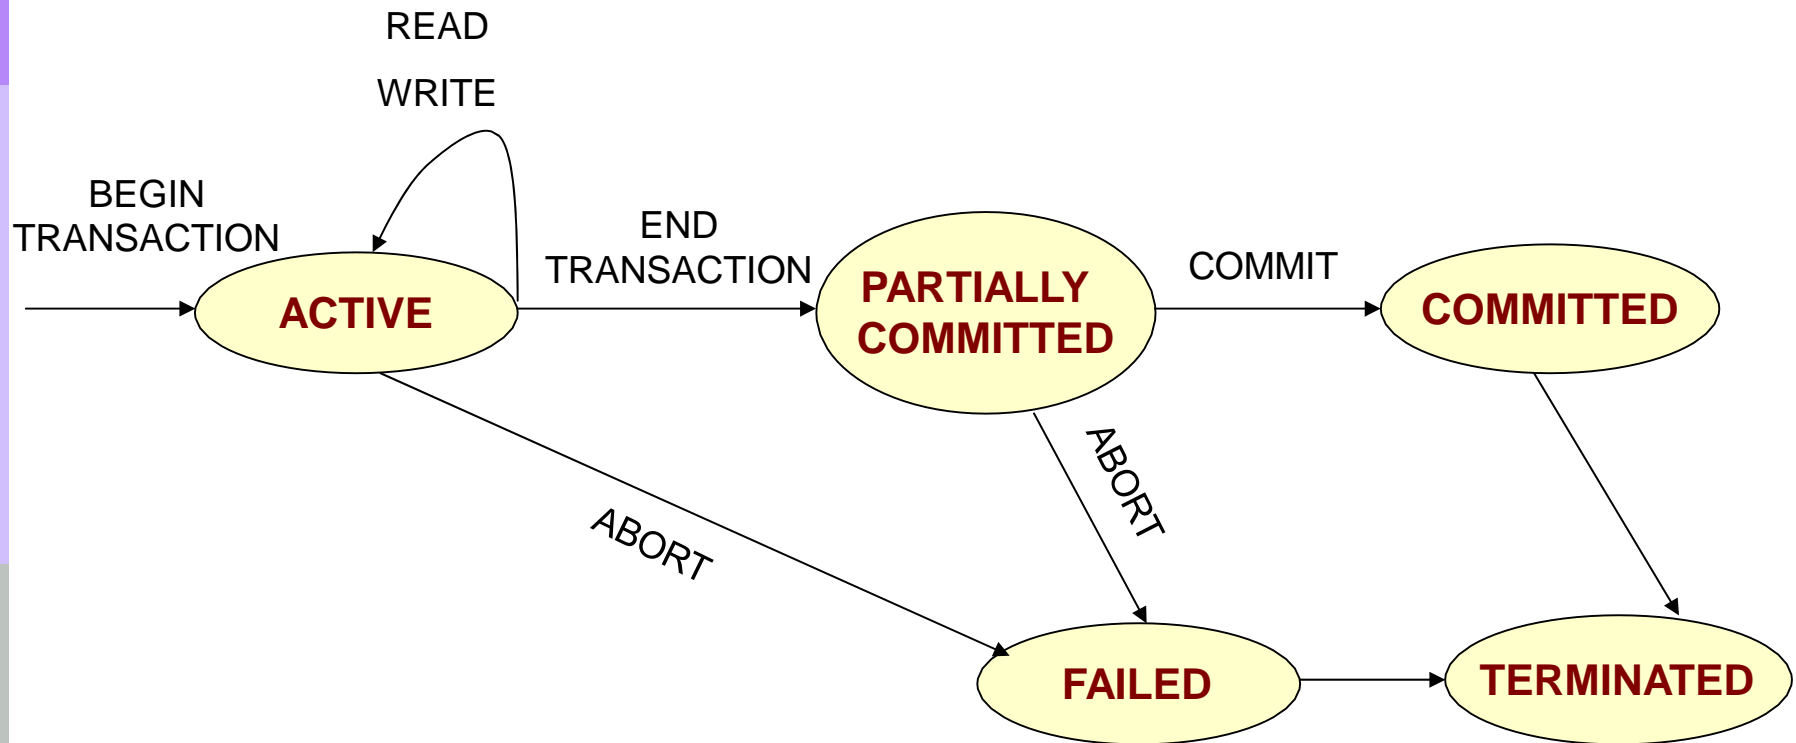
# Transaction States and Additional Operations

---

- ❖ A transaction is an **atomic unit of work that is either completed in its entirety or not done at all**
- ❖ Operations of a Transaction:
  - BEGIN TRANSACTION
  - READ OR WRITE (ACTIVE)
  - END TRANSACTION
  - COMMIT TRANSACTION / PARTIAL COMMIT
  - FAILED
  - ROLLBACK (ABORT)

# State Transition Diagram

## ❖ States of Transaction Execution



# System Log

---

- ❖ To be able to **recover from failures that affect transaction**
- ❖ System maintains a log to **keep track of all transaction operation that affects the database item**
- ❖ This is required for recovery from failure
- ❖ **Log is kept on DISK** → expected to be robust except catastrophic failure
- ❖ Log is **periodically backed up to archival storage**
- ❖ List of entries present in a log file → **LOG RECORD**

# Log Records

---

- ❖ [start\_transaction, T]
- ❖ [write\_item, T, X, old\_value, new\_value]
- ❖ [read\_item, T, X]
- ❖ [commit, T]
- ❖ [abort, T]

**UNDO**

**REDO**

# Commit Point of a Transaction

---

- ❖ A transaction T reaches its commit point when all its operation that access the database item have been executed successfully
- ❖ Effect of all the transaction operations on the database have been recorded in the LOG
- ❖ Beyond the commit point transaction is said to be committed and is permanently recorded on database

# How Log Records and Database are Affected on the Execution of any Query

---

- ❖ How Log records and data base are affected on the execution of the following query?
- ❖ Query1:  
UPDATE saving SET balance=500 where accno=1;  
COMMIT;
- ❖ Query2:  
UPDATE saving SET balance=500 where accno=1;  
ROLLBACK;

# Steps of Transaction Execution

---

**Step1:** The system will check about the **availability of the table** as well as the **attributes**

**Step2:** Whether there is any **indexing/ hashing available on search key**

**Step3:** If **index** exists search and **retrieve the data** accordingly

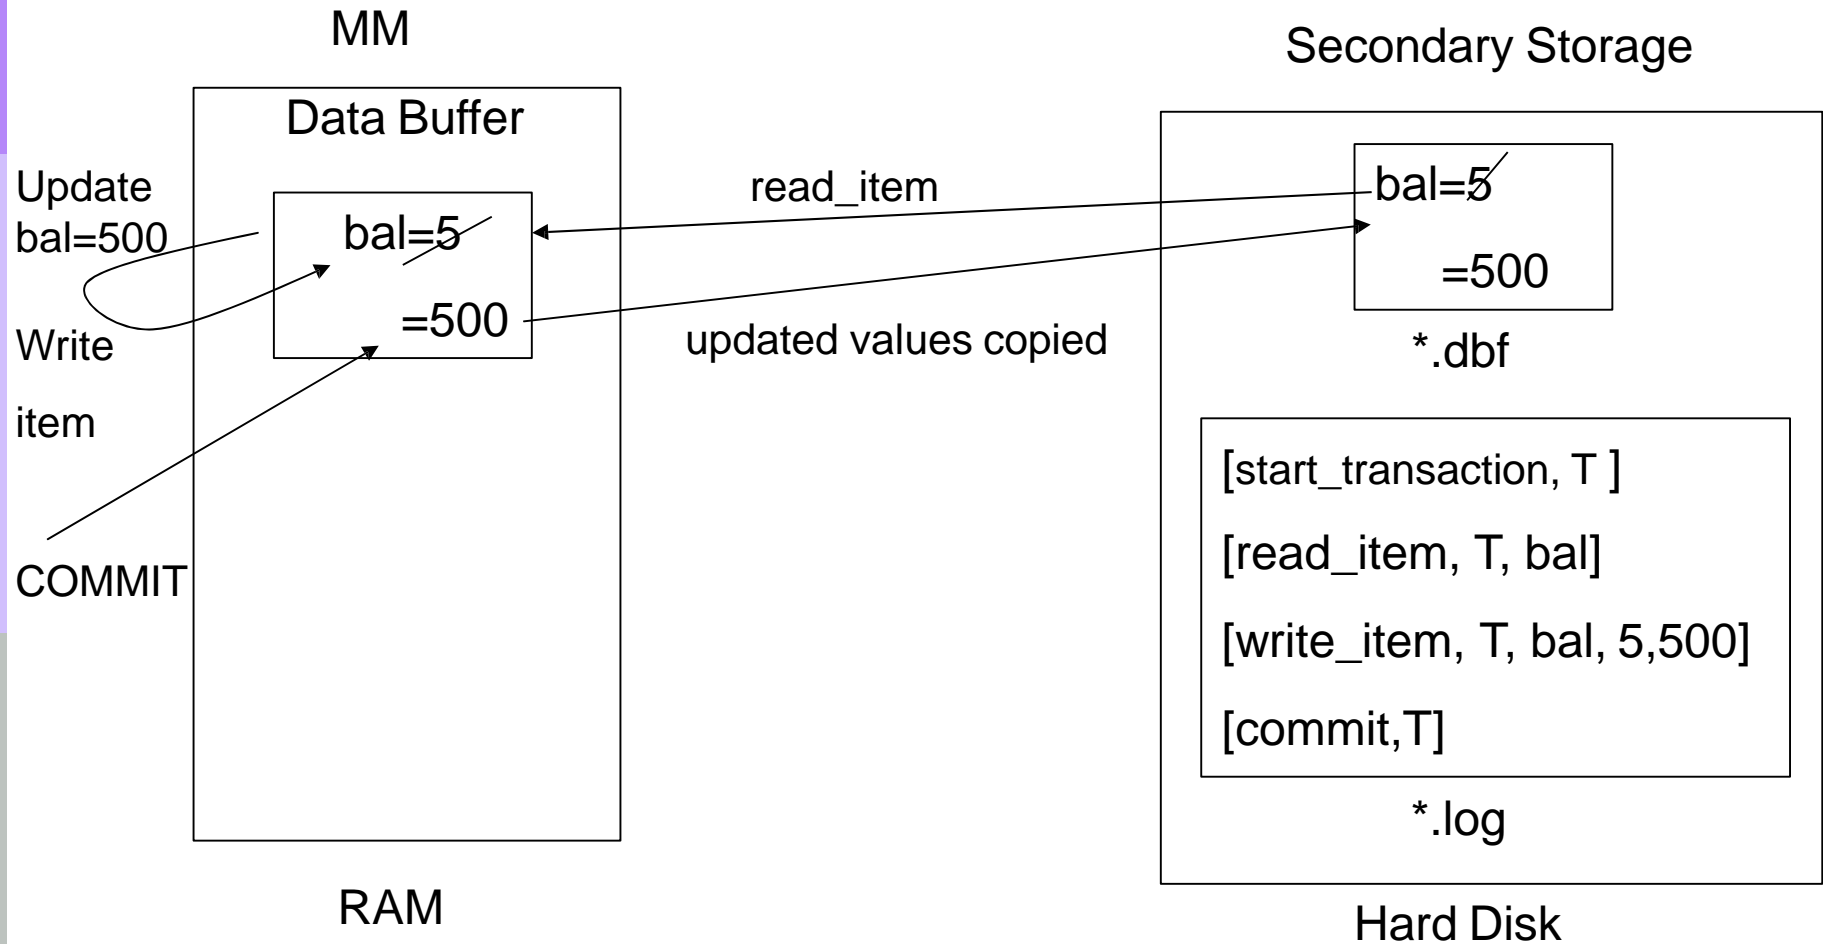
Else **retrieve the data** based on **search key values** only

**Step4:** **Update the database in MM and LOG** as shown in the figure

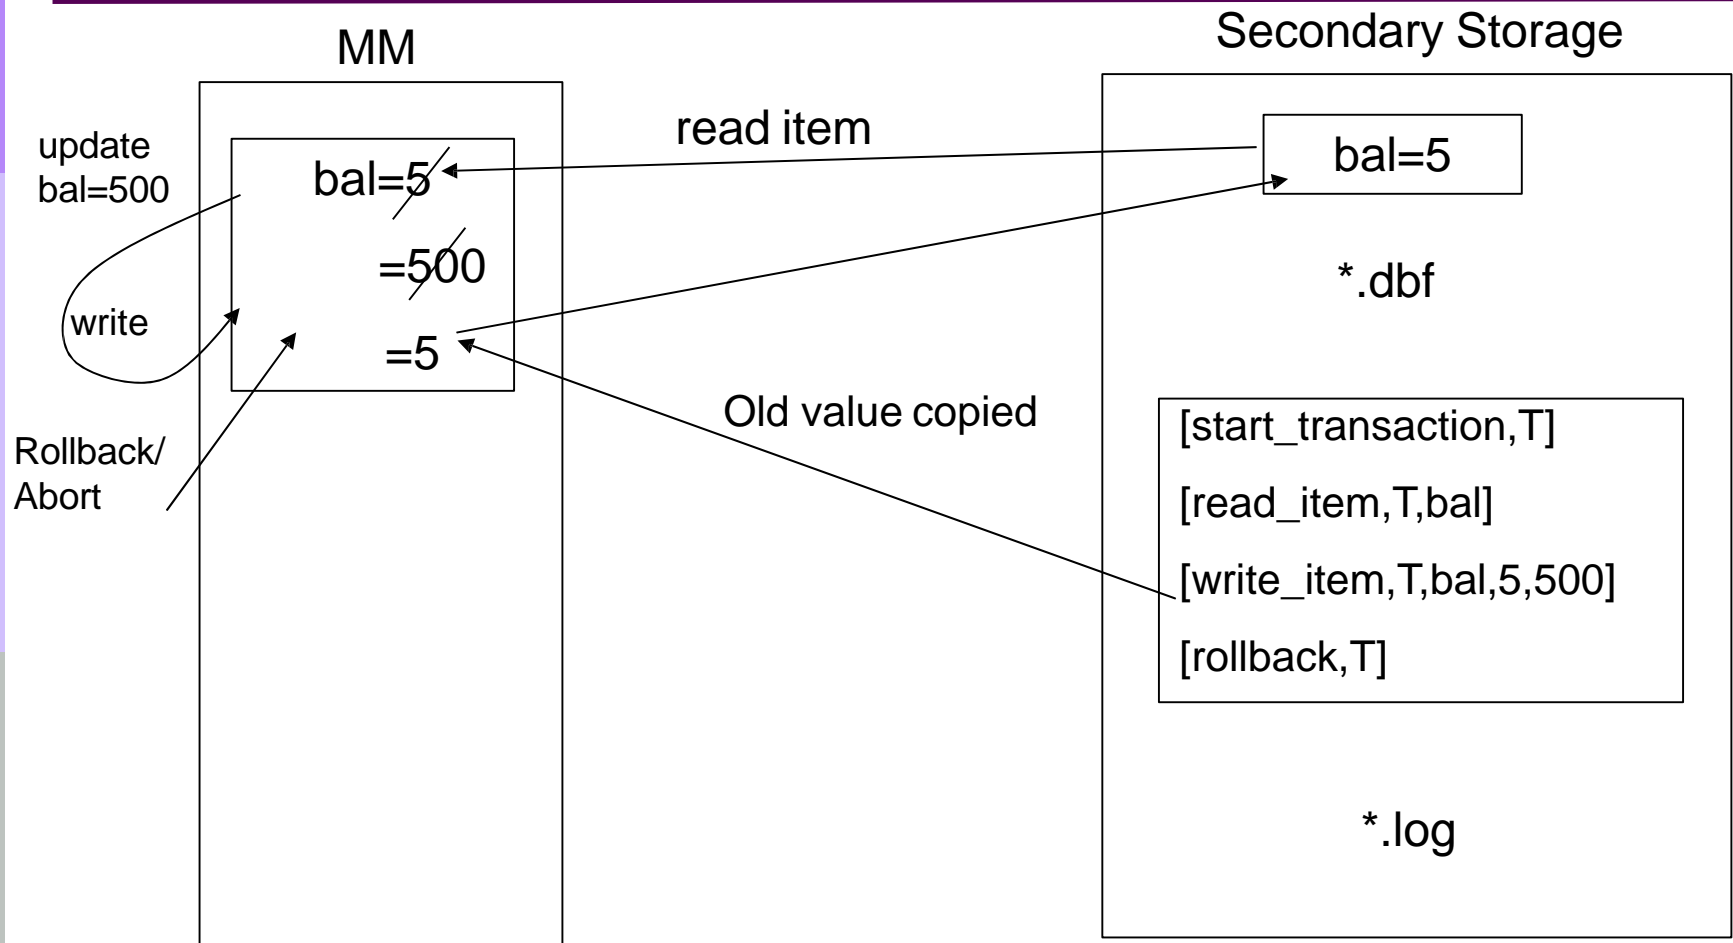
**Step4:** **COMMIT/ROLLBACK** as required



# Steps of Transaction Execution: Query1



# Steps of Transaction Execution: Query2

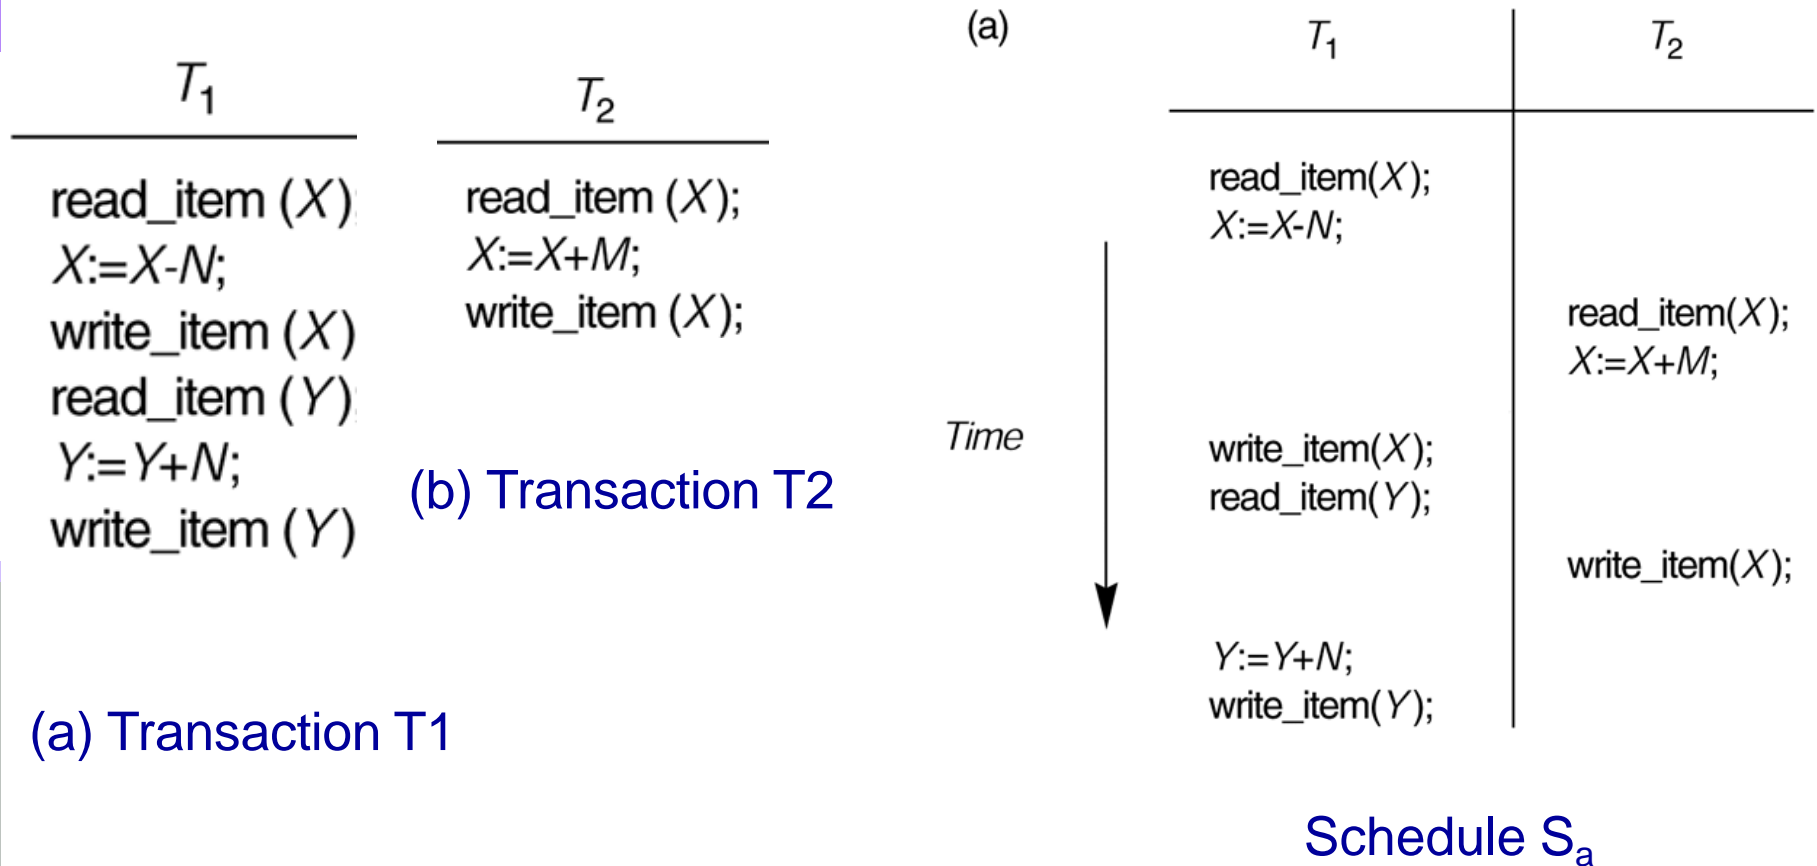


# Schedule

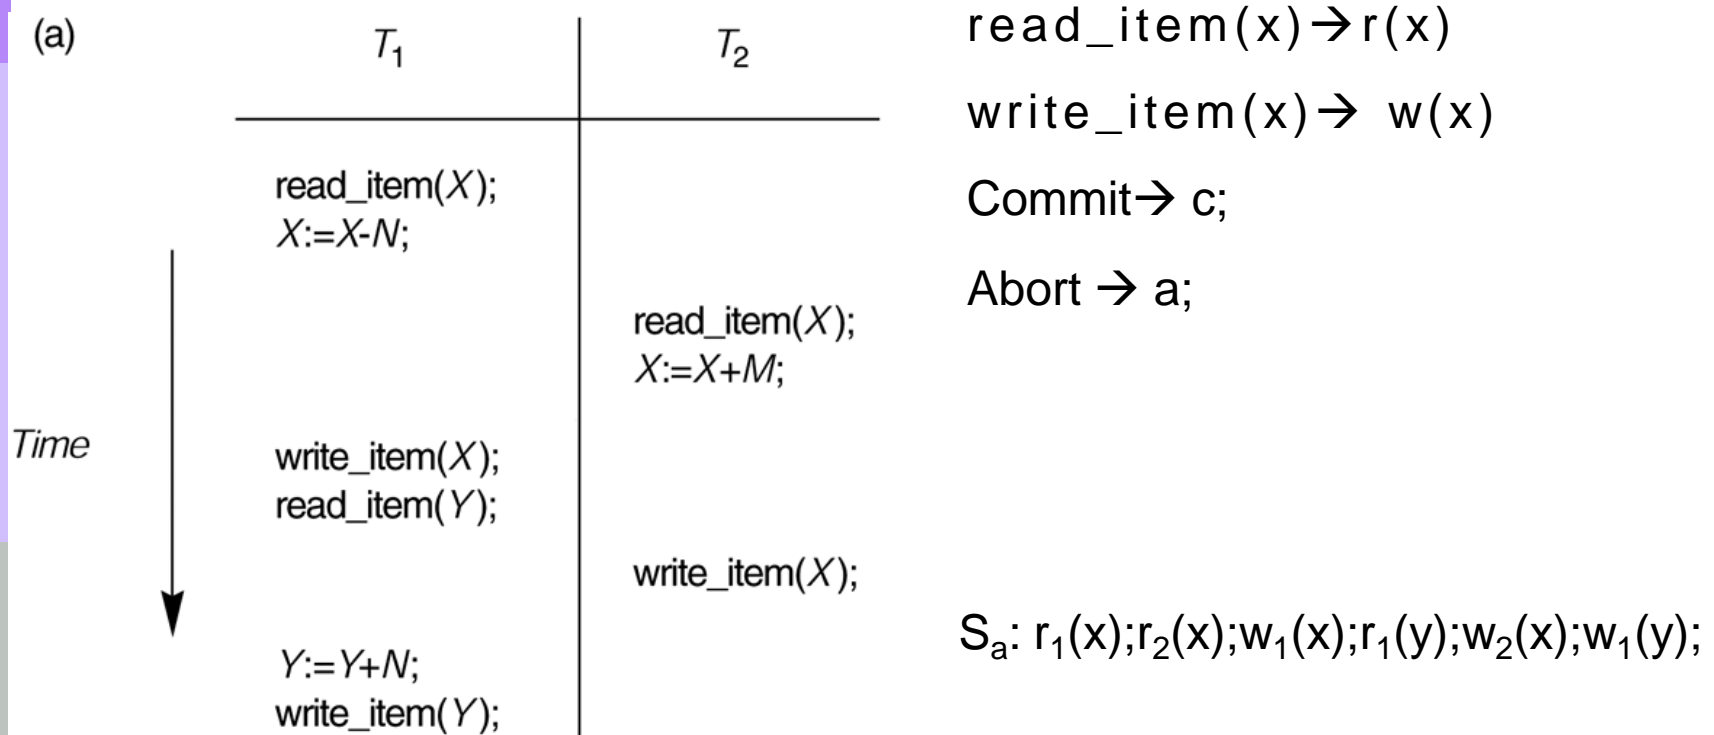
---

- ❖ When transactions are executing concurrently in an interleaved fashion then the **order of execution of operations** from the various transactions is known as a schedule
- ❖ A **schedule (history)**  $S$  of  $n$  transactions  $T_1, T_2, \dots, T_n$  is an **ordering of the operations of the transactions** subject to the constraint that, for each transaction  $T_i$  that participates in  $S$ , the operations of  $T_i$  in  $S$  must appear in the same order in which they occur in  $T_i$ .
- ❖ **Total order** Schedule Vs. **Partial Ordered** Schedule

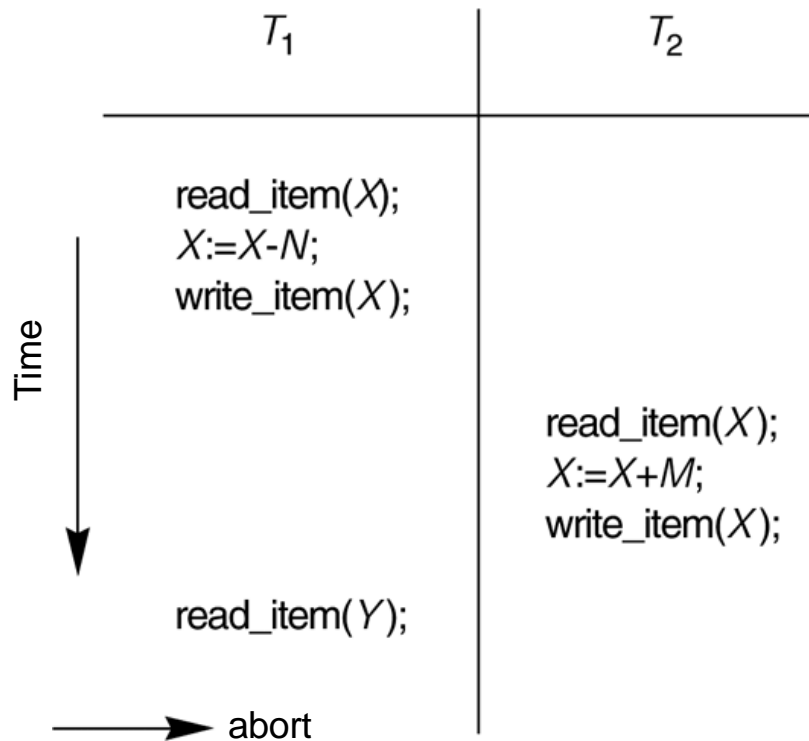
# Example of Schedule



# Example of Schedule



# Example of Schedule



$S_b: r_1(x); w_1(x); r_2(x); w_2(x); r_1(y); a_1;$

# Conflicting Operations in a Schedule

---

- ❖ Two operations in a schedule are said to **conflict** if they satisfy all the three of the following:
  - They belong to different transaction;
  - They access the same item  $x$ ;
  - At least one of them is a  $\text{write\_item}(x)$ ;

For example, in schedule  $S_a$

$S_a: r_1(x); r_2(x); w_1(x); r_1(y); w_2(x); w_1(y);$

1)  $r_1(x); w_2(x);$

2)  $r_2(x); w_1(x);$

3)  $w_1(x); w_2(x);$

# Conflicting Operations in a Schedule

---

- ❖ Test whether following operations **conflict or not?**

$S_a: r_1(x); r_2(x); w_1(x); r_1(y); w_2(x); w_1(y);$

For example, in schedule  $S_a$ :

1)  $r_1(x); r_2(x); \dots\dots\dots$

2)  $r_1(x); w_1(x); \dots\dots\dots$

3)  $w_2(x); w_1(y); \dots\dots\dots$



# Transaction Schedule

---

When transaction are executed concurrently in an interleaved fashion.

A schedule  $S$  of  $n$  transactions  $T_1, T_2, \dots, T_n$ :

- All the instructions of the participating transactions must be present in  $S$ .

- The relative ordering of the instructions of each individual transaction must be preserved.

- However, the instructions of the different transactions may be in interleaving fashion.

A transaction that successfully completes its execution will have a commit instructions as the last statement (will be omitted if it is obvious).

A transaction that fails to successfully complete its execution will have an abort instructions as the last statement (will be omitted if it is obvious).

# Schedule 1

Let  $T_1$  transfer \$50 from  $A$  to  $B$ , and  $T_2$  transfer 10% of the balance from  $A$  to  $B$ .

A serial schedule in which  $T_1$  is followed by  $T_2$ :

$T_1$	$T_2$
read( $A$ ) $A := A - 50$ write( $A$ ) read( $B$ ) $B := B + 50$ write( $B$ )	read( $A$ ) $temp := A * 0.1$ $A := A - temp$ write( $A$ ) read( $B$ ) $B := B + temp$ write( $B$ )

## Schedule 2

- A serial schedule where  $T_2$  is followed by  $T_1$
- 

$T_1$	$T_2$
<pre>read(A) A := A - 50 write(A) read(B) B := B + 50 write(B)</pre>	<pre>read(A) temp := A * 0.1 A := A - temp write(A) read(B) B := B + temp write(B)</pre>

# Schedule 3

Let  $T_1$  and  $T_2$  be the transactions defined previously. The following schedule is not a serial schedule, but it is *equivalent* to Schedule 1.

$T_1$	$T_2$
read( $A$ ) $A := A - 50$ write( $A$ )	read( $A$ ) $temp := A * 0.1$ $A := A - temp$ write( $A$ )
read( $B$ ) $B := B + 50$ write( $B$ )	read( $B$ ) $B := B + temp$ write( $B$ )

In Schedules 1, 2 and 3, the sum  $A + B$  is preserved.

# Schedule 4

The following concurrent schedule does not preserve the value of  $(A + B)$ .

$T_1$	$T_2$
read( $A$ ) $A := A - 50$	read( $A$ ) $temp := A * 0.1$ $A := A - temp$ write( $A$ ) read( $B$ )
write( $A$ ) read( $B$ ) $B := B + 50$ write( $B$ )	         $B := B + temp$ write( $B$ )

# Serializability

---

- **Basic Assumption:** Each transaction preserves database consistency.
  - ✓ Thus serial execution of a set of transactions preserves database consistency.
- **Serial schedule:** A schedule  $S$  is serial if, for every transaction  $T$  participating in the schedule, all the operations of  $T$  are executed consecutively in the schedule. Otherwise, the schedule is called nonserial schedule.
- **Serializable schedule:** A schedule  $S$  is serializable if it is equivalent to some serial schedule of the same  $n$  transactions.

# Serializability

---

- **Result equivalent:** Two schedules are called result equivalent if they produce the same final state of the database.
  - ✓ A (possibly concurrent) schedule is serializable if it is equivalent to a serial schedule. Different forms of schedule equivalence give rise to the notions of:
    1. **conflict serializability**
    2. **view serializability**
- We ignore operations other than read and write instructions, and we assume that transactions may perform arbitrary computations on data in local buffers in between reads and writes. Our simplified schedules consist of only read and write instructions.

# Conflicting Instructions

- 
- Instructions  $I_i$  and  $I_j$  of transactions  $T_i$  and  $T_j$  respectively, **conflict** if and only if there exists some item  $Q$  accessed by both  $I_i$  and  $I_j$ , and at least one of these instructions wrote  $Q$ .
    - $I_i = \text{read}(Q)$ ,  $I_j = \text{read}(Q)$ .  $I_i$  and  $I_j$  don't conflict.
    - $I_i = \text{read}(Q)$ ,  $I_j = \text{write}(Q)$ . They conflict.
    - $I_i = \text{write}(Q)$ ,  $I_j = \text{read}(Q)$ . They conflict
    - $I_i = \text{write}(Q)$ ,  $I_j = \text{write}(Q)$ . They conflict
  - Intuitively, a conflict between  $I_i$  and  $I_j$  forces a (logical) temporal order between them.
    - If  $I_i$  and  $I_j$  are consecutive in a schedule and they do not conflict, their results would remain the same even if they had been interchanged in the schedule.



# Conflicting Serializable

---

- **Conflict equivalent:** Two schedules are said to be conflict equivalent if the order of any two conflicting operations is the same in both schedules.
- **Conflict serializable:** A schedule  $S$  is said to be conflict serializable if it is conflict equivalent to some serial schedule  $S'$ .

# Conflict Serializability (Cont.)

Schedule 3 can be transformed into Schedule 6, a serial schedule where  $T_2$  follows  $T_1$ , by series of swaps of non-conflicting instructions.

Therefore Schedule 3 is conflict serializable.

$T_1$	$T_2$
read( $A$ ) write( $A$ )	read( $A$ ) write( $A$ )
read( $B$ ) write( $B$ )	
	read( $B$ ) write( $B$ )

Schedule 3

$T_1$	$T_2$
read( $A$ ) write( $A$ ) read( $B$ ) write( $B$ )	read( $A$ ) write( $A$ ) read( $B$ ) write( $B$ )

Schedule 6

# Conflict Serializability (Cont.)

---

Example of a schedule that is not conflict serializable:

$T_3$	$T_4$
read( $Q$ )	write( $Q$ )
write( $Q$ )	

We are unable to swap instructions in the above schedule to obtain either the serial schedule  $\langle T_3, T_4 \rangle$ , or the serial schedule  $\langle T_4, T_3 \rangle$ .

# View Serializability

---

- Let  $S$  and  $S'$  be two schedules with the same set of transactions.  $S$  and  $S'$  are **view equivalent** if the following three conditions are met:
1. For each data item  $Q$ , if transaction  $T_i$  reads the initial value of  $Q$  in schedule  $S$ , then transaction  $T_i$  must, in schedule  $S'$ , also read the initial value of  $Q$ .
  2. For each data item  $Q$  if transaction  $T_i$  executes **read**( $Q$ ) in schedule  $S$ , and that value was produced by transaction  $T_j$  (if any), then transaction  $T_i$  must in schedule  $S'$  also read the value of  $Q$  that was produced by transaction  $T_j$ .
  3. For each data item  $Q$ , the transaction (if any) that performs the final **write**( $Q$ ) operation in schedule  $S$  must perform the final **write**( $Q$ ) operation in schedule  $S'$ .

**As can be seen, view equivalence is also based purely on reads and writes alone.**

# View Serializability (Cont.)

- A schedule  $S$  is **view serializable** if it is view equivalent to a serial schedule.
- Every conflict serializable schedule is also view serializable.
- Conflict serializability is stricter than view serializability. With unconstrained write (or blind write), a schedule that is view serializable is not necessarily conflict serializable.
- ✓ Below is a schedule which is view-serializable but *not* conflict serializable.

$T_3$	$T_4$	$T_6$
read( $Q$ )	write( $Q$ )	
write( $Q$ )		
		write( $Q$ )

- ✓ What serial schedule is above equivalent to?
- ✓ Every view serializable schedule that is not conflict serializable has **blind** writes.

# Testing for Serializability

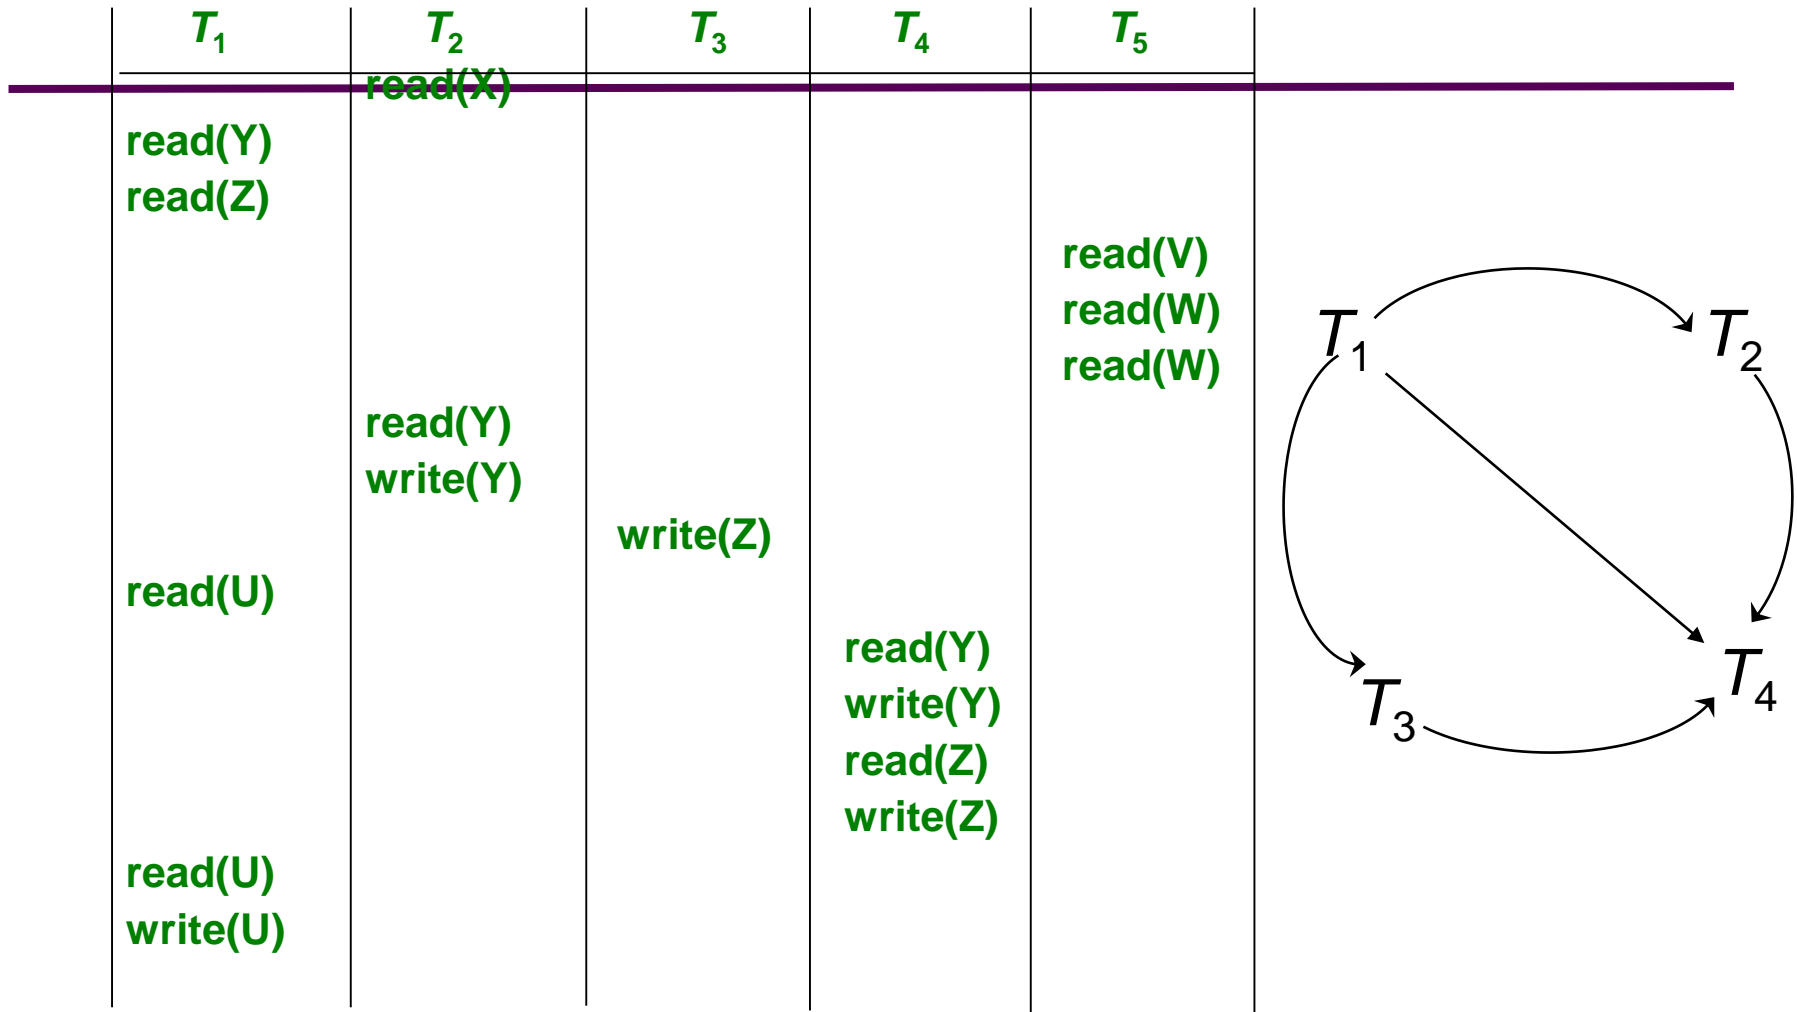
---

## Testing for conflict serializability

### Algorithm:

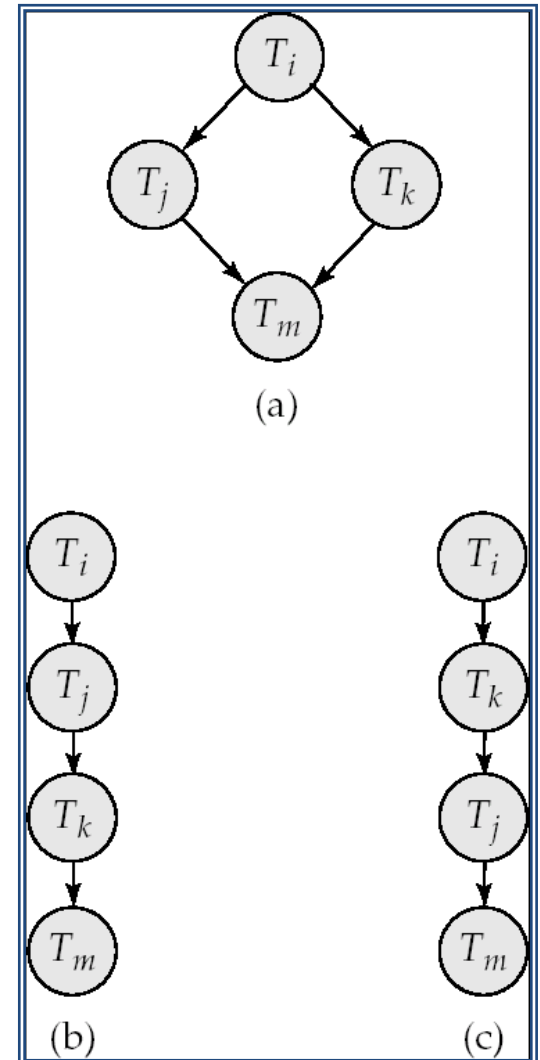
1. Looks at only read\_Item (X) and write\_Item (X) operations
2. Constructs a precedence graph (serialization graph) - a graph with directed edges
3. An edge is created from  $T_i$  to  $T_j$  if one of the operations in  $T_i$  appears before a conflicting operation in  $T_j$
4. The schedule is serializable if and only if the precedence graph has no cycles.

# Example Schedule (Schedule A) + Precedence Graph



# Test for Conflict Serializability

- A schedule is conflict serializable if and only if its precedence graph is acyclic.
- Cycle-detection algorithms exist which take order  $n^2$  time, where  $n$  is the number of vertices in the graph.  
(Better algorithms take order  $n + e$  where  $e$  is the number of edges.)
- If precedence graph is acyclic, the serializability order can be obtained by a *topological sorting* of the graph.
  - ✓ This is a linear order consistent with the partial order of the graph.
  - ✓ For example, a serializability order for Schedule A would be  
 $T_5 \rightarrow T_1 \rightarrow T_3 \rightarrow T_2 \rightarrow T_4$   
Are there others?





# Test for View Serializability

---

- The precedence graph test for conflict serializability cannot be used directly to test for view serializability.
  - ✓ Extension to test for view serializability has cost exponential in the size of the precedence graph.
- The problem of checking if a schedule is view serializable falls in the class of *NP*-complete problems.
- - ✓ Thus existence of an efficient algorithm is *extremely* unlikely.
- However practical algorithms that just check some sufficient conditions for view serializability can still be used.

# Recoverable Schedules

---

Need to address the effect of transaction failures on concurrently running transactions.

- **Recoverable schedule** — if a transaction  $T_j$  reads a data item previously written by a transaction  $T_i$ , then the commit operation of  $T_i$  appears before the commit operation of  $T_j$ .
- ✓ The following schedule (Schedule 11) is not recoverable if  $T_9$  commits immediately after the read

$T_8$	$T_9$
read(A) write(A)  read(B)	  read(A)

- ✓ If  $T_8$  should abort,  $T_9$  would have read (and possibly shown to the user) an inconsistent database state. Hence, database must ensure that schedules are recoverable.

# Cascading Rollbacks

---

**Cascading rollback** – a single transaction failure leads to a series of transaction rollbacks. Consider the following schedule where none of the transactions has yet committed (so the schedule is recoverable)

$T_{10}$	$T_{11}$	$T_{12}$
read( $A$ ) read( $B$ ) write( $A$ )	read( $A$ ) write( $A$ )	read( $A$ )

If  $T_{10}$  fails,  $T_{11}$  and  $T_{12}$  must also be rolled back.  
Can lead to the undoing of a significant amount of work

# Cascadeless Schedules

---

**Cascadeless schedules** — cascading rollbacks cannot occur; for each pair of transactions  $T_i$  and  $T_j$  such that  $T_j$  reads a data item previously written by  $T_i$ , the commit operation of  $T_i$  appears before the read operation of  $T_j$ . Every cascadeless schedule is also recoverable  
It is desirable to restrict the schedules to those that are cascadeless

# Concurrency Control

---

- A database must provide a mechanism that will ensure that all possible schedules are
  - ✓ either conflict or view serializable, and
  - ✓ are recoverable and preferably cascadeless
- A policy in which only one transaction can execute at a time generates serial schedules, but provides a poor degree of concurrency
  - ✓ Are serial schedules recoverable/cascadeless?
- Testing a schedule for serializability *after* it has executed is a little too late!
- **Goal** – to develop concurrency control protocols that will assure serializability.

# Concurrency Control vs. Serializability Tests

---

- Concurrency-control protocols allow concurrent schedules, but ensure that the schedules are conflict/view serializable, and are recoverable and cascadeless.
- Concurrency control protocols generally do not examine the precedence graph as it is being created
  - ✓ Instead a protocol imposes a discipline that avoids nonserializable schedules.
- Different concurrency control protocols provide different tradeoffs between the amount of concurrency they allow and the amount of overhead that they incur.
- Tests for serializability help us understand why a concurrency control protocol is correct.

# Complete Schedule

---

- ❖ A **schedule**  $S$  of  $n$  transactions  $T_1, T_2, \dots, T_n$  is said to be a complete schedule if the following conditions hold:
  - The operations in  $S$  are exactly those operations in  $T_1, T_2, \dots, T_n$ , including a commit or abort operation as the last operation for each transaction in the schedule
  - For any pair of operations from the same transaction  $T_i$ , their order of appearance in  $S$  is the same as their order of appearance in  $T_i$
  - For any two conflicting operations, one of the two must occur before the other in the schedule

---

# Thank You

# ?